Report A · X Account Intelligence & Algorithm Audit · generated 2026-10-23 08:58 UTC

Nicolas Finet (@nifinet) — Complete X Account Audit

Window 2026-04-26 → 2026-10-22 (180 days; 90-day figures marked) · public data only · algorithm repository xai-org/x-algorithm@07d159b

> How to read this report: every finding carries an evidence class — repo verified · account observed · statistical association · hypothesis · unsupported · insufficient. Associations are not causes; the published repository does not describe every live production system; no metric here is an X algorithm score.

1. Executive summary

FOLLOWERS	MEDIAN VIEWS / AUTHORED POST	TREND PER 30 DAYS	MEDIAN VIEWS / COMMENT
11,613	**1,258**	**0.92x**	**47**
snapshot 2026-10-23	n=157 · 180d	90% CI 0.82–1.03	n=226 conversational

Top 5 potential barriers to organic growth

#	Issue	Status	Confidence	Evidence class	Priority
1	Content format drives most of the reach difference	supported	high	statistical association	0.89
2	Low conversation generation per impression	supported	medium	statistical association	0.68
3	Off-core topics underperform core topics	supported	medium	statistical association	0.67
4	Reach depends on rare breakouts	supported	high	account observed	0.65
5	Commenting is low-depth and concentrated on a small circle	supported	medium	account observed	0.64

Priority = 0.30·evidence strength + 0.25·likely impact + 0.20·business relevance + 0.15·confidence + 0.10·ease of testing (transparent internal weighting).

Stop / continue / start

Stop

- Reduce: Length: short (15-49) (0.18x median, n=24)
- Reduce: Educational value (0-3): 0.0 (0.19x median, n=34)
- Reduce: Time of day (UTC): 06-11 (0.20x median, n=21)
- Reduce: Narrative structure: other (0.21x median, n=28)
- Reduce: Opening hook type: other (0.22x median, n=24)
- Stop publishing a 3rd authored post on the same day
- Stop generic/emoji-only comments; they earn little exposure
- Stop mass identical public 'Sent' replies (test DM/pinned delivery)

Continue

- Keep: Content format: x_article (15.12x median, n=30)
- Keep: Opening hook type: how_to (8.87x median, n=33)
- Keep: Educational value (0-3): 2.0 (6.97x median, n=46)
- Keep: Opening hook type: story (6.24x median, n=8)
- Keep: Narrative structure: tutorial (6.10x median, n=46)

Start

- Start measuring: daily metric snapshots + follower history (monitoring)
- Start the randomised hook and closing-prompt tests (see 30-day plan)
- Start substantive comments on larger in-niche accounts within their first hour
- Start replying to early commenters within 60 minutes to build conversation depth

2. Data completeness and confidence

Data	Status	Detail
Posts, replies, quotes, reposts (180 days)	available	632 records from the ThoughtPilot console store
Views (public view count)	available	618 of 632 records have views
Likes, replies, reposts, quotes, bookmarks	available	public counts
Metric scrape time	partial	the store does not record when metrics were scraped; observations are stamped at import time. Spot-check vs an August scrape: stored views were equal or higher for overlapping posts (metrics are refreshed, not frozen).
Engagement over time per post	unavailable	one observation per post; monitoring (Phase 2) builds the time series
Follower count	partial	single snapshot on 2026-10-23; no historical follower growth
Profile bio / account creation date	unavailable	user-info returned no description or creation date
Media type (image/video)	unavailable	not in source records; t.co URLs may be links or media
Profile visits, follows, link clicks, audience demographics	unavailable	private analytics; this audit uses public data only (by decision)
Reply parent author	inferred	from leading @mention for 385 replies; confirmed by tweet-detail for 31
Reply parent post (time, text, engagement)	sample	tweet-detail collected for 71 replies across clients (most recent conversational comments); 230 conversational comments in total for this account
Who engaged with the account's posts	available	392 public replies/quotes by others (client-engagers); likes are not exposed
Content/reply labels	available	label methods: {'llm:claude-sonnet-session-review': 577, 'none': 55}
Immature posts (<72h old at observation)	excluded from baselines	8 posts

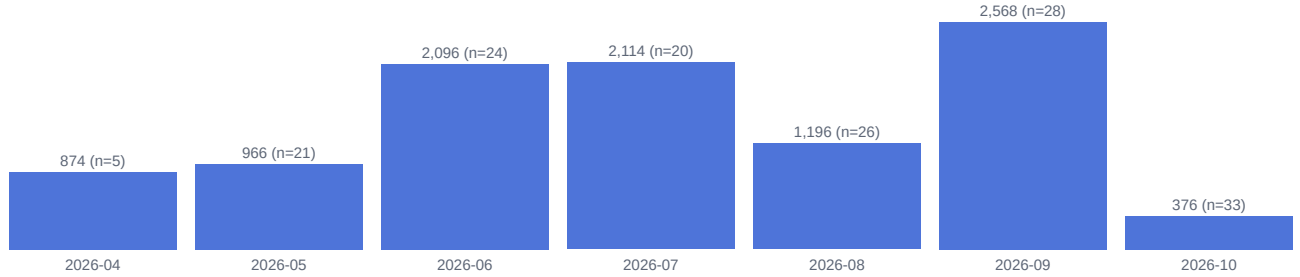
3. Historical performance

Reach by post type (mature posts, 180 days)

Segment	n	Median	p25	p75	p90	Mean	Median / follower
authored (original + quote)	157	1,258	407	4,214	17,201	73,893	10.83%
original	91	1,931	514	6,538	25,145	125,492	16.63%
quote	66	754	281	2,627	4,741	2,750	6.49%
reply to others	411	25	9	64	172	1,323	0.22%
self-reply (thread)	18	140	79	183	439	174	1.21%

90-day window: median authored views 1,183 (n=91). The top 5% of authored posts produced 96% of authored views (top single post 85%), so medians, not totals, describe a typical post. Relative to each post's trailing 28-day baseline: 36% exceeded 2x, 42% fell below 0.5x.

Monthly median views (authored) and posting volume



Month	Authored posts	Median views	Total views	Replies to others	Median reply views
2026-04	5	874	6,179	0	–
2026-05	21	966	44,120	2	184
2026-06	24	2,096	367,676	21	19
2026-07	20	2,114	930,245	68	62
2026-08	26	1,196	66,884	43	54
2026-09	28	2,568	239,087	214	10
2026-10	33	376	9,947,024	64	40

Engagement efficiency (per 1,000 views; observed outcomes, not ranking scores)

Type	likes /1k (median)	replies /1k (median)	reposts /1k (median)	quotes /1k (median)	bookmarks /1k (median)	eng /1k (median)
authored	4.99	0.37	0.38	0.00	2.44	10.88
reply to others	0.00	0.00	0.00	0.00	0.00	0.00

Unavailable: profile visits per 1k (private analytics (account owner export) — not collected); follows per 1k (private analytics — not collected); link clicks per 1k (private analytics — not collected); video watch metrics (private analytics; media type also not present in source data); dwell time (not exposed publicly at all).

4. Posting audit

Characteristics associated with higher or lower reach

Each level vs all other authored posts. Ratio of medians with 90% bootstrap CI; q = Benjamini–Hochberg adjusted p across 82 comparisons. Only q < 0.1 shown as signals. Associations only.

Feature	Level	n	Median views	Ratio	90% CI	q	Signal
Content format	quote_post	66	754	0.39x	0.21–0.93	0.004	lower
Content format	short_post	14	398	0.29x	0.18–0.49	0.092	lower
Content format	x_article	30	11,630	15.12x	8.28–29.67	0.000	higher
Topic	AI agents & GTM automation	33	4,175	4.18x	2.33–8.52	0.001	higher
Topic	AI industry & tools news	12	350	0.27x	0.14–1.42	0.081	lower
Topic	buying signals & intent data	45	2,381	2.75x	1.84–5.09	0.006	higher
Topic	other	15	378	0.26x	0.15–0.37	0.001	lower
Topic	personal & lifestyle	9	326	0.24x	0.15–0.46	0.092	lower
Opening hook type	announcement	8	4,488	3.71x	1.84–6.72	0.089	higher
Opening hook type	how_to	33	7,667	8.87x	4.82–18.49	0.000	higher
Opening hook type	observation	37	874	0.61x	0.24–1.07	0.077	lower
Opening hook type	other	24	382	0.22x	0.15–0.34	0.001	lower
Opening hook type	story	8	7,547	6.24x	2.44–17.66	0.022	higher
Narrative structure	opinion	29	612	0.35x	0.17–0.66	0.005	lower
Narrative structure	other	28	398	0.21x	0.15–0.32	0.000	lower
Narrative structure	promotion	10	4,488	3.79x	2.48–7.18	0.024	higher
Narrative structure	tutorial	46	4,690	6.10x	3.07–11.46	0.000	higher
Length	short (15-49)	24	312	0.18x	0.13–0.30	0.000	lower
Length	very short (<15 words)	56	2,336	1.97x	0.70–6.52	0.067	higher
Contains t.co URL (link or media)	False	93	812	0.27x	0.14–0.48	0.001	lower
Contains t.co URL (link or media)	True	64	3,019	3.72x	2.03–7.06	0.001	higher
Opinion strength (0-3)	0.0	76	2,336	2.73x	1.48–5.26	0.009	higher
Opinion strength (0-3)	2.0	34	626	0.35x	0.25–0.68	0.040	lower
Educational value (0-3)	0.0	34	398	0.19x	0.15–0.30	0.000	lower
Educational value (0-3)	1.0	33	612	0.45x	0.19–1.07	0.040	lower
Educational value (0-3)	2.0	46	4,462	6.97x	3.02–10.54	0.000	higher
Weekday	Monday	30	2,784	2.43x	1.26–6.25	0.077	higher
Weekday	Saturday	12	386	0.27x	0.13–0.39	0.002	lower
Time of day (UTC)	06-11	21	376	0.20x	0.14–0.29	0.000	lower
Time of day (UTC)	12-17	68	1,762	1.91x	1.07–3.63	0.016	higher

Multivariable model (controls: month fixed effects; HC3 robust errors)

log_views ~ C(format) + C(month) + C(topic) + tco_url_excl_articles + has_list + uses_data + personal_experience + promotional + keyword_cta + log_words · n=157 · R^2 0.55. Multipliers are relative to the reference level of each factor.

Term	Multiplier	90% CI	p
C(format)[T.x_article]	132.22x	29.17–599.31	0.000
has_list	1.96x	1.20–3.19	0.024
log_words	1.47x	1.11–1.94	0.026
C(format)[T.quote_post]	1.68x	1.03–2.74	0.083
C(topic)[T.other/rare]	0.50x	0.26–0.98	0.092
promotional	1.88x	0.94–3.76	0.133
C(topic)[T.outbound & cold email]	0.47x	0.20–1.11	0.149
C(format)[T.short_post]	2.50x	0.54–11.46	0.323
uses_data	1.22x	0.79–1.87	0.447
C(topic)[T.buying signals & intent data]	0.79x	0.47–1.33	0.462
C(topic)[T.personal & lifestyle]	2.53x	0.31–20.84	0.470
C(topic)[T.agency business & scaling]	0.61x	0.18–2.10	0.511
personal_experience	0.86x	0.47–1.55	0.667
C(format)[T.thread]	0.81x	0.35–1.85	0.672

Time and frequency

6.3 authored posts and 16.2 replies per week; active on 82% of days; up to 97 replies in a single day.

Same-day order	n	Median vs baseline
1st post of the day	117	1.10x
2nd post of the day	26	1.38x
3rd+ post of the day	9	0.23x

Gap since previous post	n	Median vs baseline
<2h	5	0.19x
2-6h	20	0.28x
6-24h	68	0.60x
>24h	58	1.70x

Reply volume vs authored reach (Spearman): same day ρ=-0.14 (p=0.075); previous 3 days ρ=-0.16 (p=0.052).

5. Commenting audit

Reply roles: conversation 230 · fulfilment 186 · self thread 18. Fulfilment replies are inferred from their template text. Each one answers a person who commented a keyword, so their volume is a lower bound on keyword-CTA comments received.

Category	n	Median views	Median likes	Share with a reply	Median effectiveness
humor	32	62	1.0	38%	23.0
low value repetitive	29	40	0.0	10%	30.7
generic agreement	28	61	1.0	39%	33.9
promotional	28	63	0.0	4%	39.5
original insight	26	38	1.0	31%	65.4
relationship building	25	54	0.0	24%	14.2
question	20	90	0.0	55%	70.6
personal experience	19	38	1.0	37%	52.7
expert explanation	16	26	0.0	31%	64.5
constructive disagreement	3	117	0.0	100%	80.3

Commenting Effectiveness Score — internal diagnostic, not an X algorithm score. Weights: relevance 0.25, originality 0.25, observed_engagement 0.2, conversation 0.15, timing 0.15 (timing unavailable → redistributed).

Thread-context sample (n=27): median delay after the parent post 77 min; 44% within the first hour; parent author replied to 15%.

Commenting questions

Are the accounts commenting on relevant posts?
53% of conversational comments are on ICP-relevant subjects (judged from reply text; parent post unavailable). account_observed

Are replies early enough to participate in active discussions?
In a sample of 27 recent comments with thread context, the median delay after the parent post was 77 minutes; 44% were within the first hour. account_observed

Do parent authors respond?
15% of sampled comments got a visible reply from the parent author. account_observed

Are comments substantive or repetitive?
37% substantive (insight, expertise, question, disagreement, experience); see category mix. account_observed

Are replies attracting new people or the same small group?
111 distinct accounts commented on; the top 5 receive 35% of comments. Who engages *with* the comments is not observable without engager collection. account_observed

Is reply activity associated with follower growth?
Insufficient data: no follower-count history exists yet (monitoring will create it). insufficient

Typical exposure per comment
median 47 views per conversational comment (n=226). account_observed

Most-viewed comments

Date	To	Comment	Category	Views
2026-10-02	@bjorndunkel	@bjorndunkel @juanjovn haha not me	humor	509,948
2026-10-11	@jjen_abel	@jjen_abel Enterprise sales: 6 months of procurement. Also enterprise sales: "you two should meet" by text.	humor	2,225
2026-08-29	@patrickc	@patrickc @paulg We purchased a German company: I can tell you that it is a nightmare administratively. Maintaining a company, closing it, dealing with employee	personal experience	1,771
2026-07-20	@jeremy_D_C	@jeremy_D_C Combien de MRR?	question	1,599
2026-06-18	@yasser_elsaid_	@yasser_elsaid_ @levelsio Ok Yasser... but did you eat steaks and lifted weight? 😆	humor	1,570
2026-07-20	@MichLieben	@MichLieben guilty 😎	humor	1,061
2026-09-12	@iamneubert	@iamneubert Love that! What did you use?	question	963
2026-08-08	@levelsio	@levelsio Fine with ChatGPT 😁 https://t.co/rHp2g7fCrM	personal experience	929

6. Audience and network

111 distinct accounts received conversational comments; the top 5 received 35% (HHI 0.037); 14 accounts received 3+ comments.

Account commented on	Comments	Median views	Total views	Share with a reply	Peer/client
@MichLieben	29	57	3,572	28%	yes
@itsalexvacca	23	45	1,332	30%	
@nikolak47	12	44	910	8%	
@Tycho_luijten	9	44	388	44%	
@NicholasDulait	6	34	254	33%	
@antinertia	4	223	922	0%	
@512banque	4	263	952	25%	
@suganthan	4	39	308	25%	
@VukConfidential	4	22	105	25%	
@Aknaimi_	3	74	193	67%	
@Kenny_GTM	3	37	128	0%	
@levelsio	3	26	974	0%	
@dagorenouf	3	144	322	33%	
@vandegans	3	11	34	33%	
@YannDecoopman	2	195	390	50%	

392 public replies/quotes from 293 distinct accounts (2026-07-27 → 2026-10-16); 89% engaged only once; the top 10 accounts made 19% of engagements; 45 engagers are also accounts @nifinet comments on (33% of engagements). Public replies/quotes that mention the account (since 2026-07-25), from the console store; includes replies to the account's comments elsewhere. Likes and reposts are not exposed.

Accounts that replied to or quoted @nifinet

Account	Replies/quotes
@itsalexvacca	23
@MichLieben	13
@nikolak47	6
@itsivanfalco	6
@Tycho_luijten	6
@VukConfidential	5
@Kenny_GTM	5
@vandegans	4
@mailankhong	4
@heyreach_io	4
@Vick_Mucunski	4
@NicholasDulait	4
@Frontal_AI	4
@suganthan	3
@boardyai	3

> Private social-graph relationships (who follows whom, mutual follows) are not inferred from public replies.

7. Potential reach barriers — full diagnostics

Content format drives most of the reach difference statistical association · supported · confidence: high

Observation	X Article posts reach 15.12x the median views of the account's other authored posts (n=30, 90% CI 8.28-29.67). Lower formats: short_post 0.29x (n=14), quote_post 0.39x (n=66).
Evidence	X Article median / other authored median: **15.12** · X Article posts: **30** · ratio CI: **8.28-29.67** · BH q-value: **0** · short_post median ratio: **0.29** · short_post posts: **14** · quote_post median ratio: **0.39** · quote_post posts: **66** · regression multiplier (month+topic controls): **132.2** · top X Article: 672224 views · top X Article: 155488 views · top X Article: 145962 views
Algorithm connection	The value model rewards predicted dwell, click-dwell and sharing (dwell, cont_click_dwell_time, share_via_copy_link weights). Long-form articles plausibly raise those predictions; the repo does not contain an article-specific boost. B05 · Published weights: shares (copy link 20), dwell, quote (home-mixer/params/param.rs:381-395) B06 · Published weights: follow_author, unexplored, continuous dwell (home-mixer/params/param.rs:396-437) B01 · Predicted heads consumed by the value model (xai-value-model/phoenix_scores.rs:1-31)
Alternative explanations	Article topics may simply be the account's strongest topics; X may surface Articles through surfaces outside the For You ranking studied here; Articles are often promoted harder (pinned, quoted by co-founders/peers)
Impact	Largest observed lever: shifting mix toward the best format changes the median outcome, not just outliers.
Recommendation	Make X Articles (with a strong native preview post) the anchor format, and turn short standalone posts into either article teasers or replies.
Validation	A/B over 4 weeks: alternate article vs non-article on matched topics; compare median views per format.
Success criteria	Article posts keep >= the current ratio vs other formats while posting volume is held constant.
Confounders	topic, promotion effort, author network resharing

Low conversation generation per impression

statistical association · supported · confidence: medium

Observation	Replies per 1,000 views: 0.93 vs peer median 1.83 (0.51x). Likes per 1,000 views: 4.16 vs 11.84. Bookmarks per 1,000 views: 2.14 vs 3.86.
Evidence	client replies per 1k views: **0.93** · peer median replies per 1k views: **1.83** · client/peer replies per 1k views: **0.51** · client engagements per 1k views: **9.25** · peer median engagements per 1k views: **22.73** · client/peer engagements per 1k views: **0.41** · client bookmarks per 1k views: **2.14** · peer median bookmarks per 1k views: **3.86** · client/peer bookmarks per 1k views: **0.55** · client likes per 1k views: **4.16** · peer median likes per 1k views: **11.84** · client/peer likes per 1k views: **0.35**
Algorithm connection	Published defaults weight a viewer's predicted reply probability at 5.0 versus 0.66 for a like; for original posts from an author the viewer mutually follows, the reply weight is raised by a further 15.0. Weights multiply the model's predicted probabilities per viewer, not raw counts, so observed reply rates are only an outcome proxy for what the model predicts. B03 · Published weights: like, reply, bidirectional boost, repost (home-mixer/params/param.rs:345-359) B08 · Weights multiply predicted probabilities; only Home Timeline engagement counts (home-mixer/params/param.rs:316-344) B15 · Bidirectional boost eligibility: mutual-follow original posts only (xai-value-model/weights.rs:84-100)
Alternative explanations	Audience composition differs from peers (more lurkers / buyers who do not reply); Lead-magnet CTAs convert replies into keyword comments counted elsewhere; Peer accounts differ in size; per-view rates partly normalise this but not fully
Impact	Conversation is the most heavily weighted positive action family; raising it is the most algorithm-aligned lever available from content alone.
Recommendation	End posts with a specific, answerable prompt tied to the post's claim; reply to early commenters within the first hour to deepen threads.
Validation	Compare replies per 1k views on posts with vs without a specific closing prompt over 4 weeks.
Success criteria	Median replies per 1k views on authored posts moves toward the peer median.
Confounders	audience size, topic, CTA type

Off-core topics underperform core topics

statistical association · supported · confidence: medium

Observation	The top two topics account for 50% of authored posts. Topics associated with higher reach: AI agents & GTM automation 4.18x, buying signals & intent data 2.75x. Lower: AI industry & tools news 0.27x, other 0.26x, personal & lifestyle 0.24x.
Evidence	share of authored posts in top-2 topics: **0.5** · AI agents & GTM automation (n=33) median ratio: **4.18** · buying signals & intent data (n=45) median ratio: **2.75** · AI industry & tools news (n=12) median ratio: **0.27** · other (n=15) median ratio: **0.26** · personal & lifestyle (n=9) median ratio: **0.24**
Algorithm connection	SimClusters and Phoenix retrieval match posts to viewers whose recent engagement resembles the post. Consistent topics make it more likely the right viewers already have related engagement history; the repo exposes no per-account topic score. A11 · SimClusters source constants (score floor, seed age, result caps) (home-mixer/sources/simclusters_source.rs:25-36) A12 · SimClusters seeds come from the viewer's recent engagements (home-mixer/sources/simclusters_source.rs:222-252) A05 · Phoenix retrieval is a two-tower model over engagement history (phoenix/README.md:99-121)
Alternative explanations	Topic effects overlap with format (tutorial topics are more often articles); Off-core posts may be personal/timely and judged on other goals
Impact	Concentrating on the strongest clusters improves the median post without adding volume.
Recommendation	Prioritise the higher-reach topics listed; reduce off-core posts or move them to replies.
Validation	Re-run the topic comparison after 6 weeks with the regression (format + month controls).
Success criteria	Share of authored posts in the higher-reach topics rises with no fall in their median.
Confounders	format, timing, promotion

Reach depends on rare breakouts account observed · supported · confidence: high

Observation	The top 5% of authored posts produced 96% of authored views; the single top post 85%. 42% of posts land below half the trailing baseline and 36% above twice it.
Evidence	share of authored views from top 5% of posts: **0.96** · share from the single top post: **0.85** · posts below 0.5x baseline: **0.42** · posts above 2x baseline: **0.36** · top post: 9900561 views · top post: 672224 views · top post: 155488 views
Algorithm connection	Out-of-network discovery starts small (cold-start pool for posts with few likes/views; retrieval-index admission triggered by likes) and only some posts escape it. Heavy-tailed outcomes are expected from this design. A07 · Cold-start retrieval pool: low-fav, low-view, fresh posts (phoenix/xrex/data/cold_pool_filter.py:16-51) A08 · Phoenix retrieval index admission is triggered by favorite events (phoenix-rankall-strato/columns/phoenix_rank_all/phoenixRankAllCandidateProcessor.strato:469-530) D19 · Cold-start lift for low-impression posts from smaller authors (home-mixer/scorers/author_cold_start.rs:188-193)
Alternative explanations	Viral posts may be driven by external resharing (large accounts, newsletters)
Impact	Averages mislead: judge experiments on medians and share-above-baseline, not totals.
Recommendation	Reverse-engineer the top posts' structure into a repeatable template; judge tests on medians.
Validation	Share of posts above 2x baseline per month.
Success criteria	Share above 2x baseline rises without lowering the median.
Confounders	external amplification

Commenting is low-depth and concentrated on a small circle account observed · supported · confidence: medium

Observation	226 conversational comments in the window; 37% are substantive and 39% are generic, repetitive or humour. The top 5 accounts receive 35% of comments; median exposure is 47.0 views per comment.
Evidence	conversational comments (mature): **226** · median views per comment: **47** · substantive share: **0.37** · generic/repetitive/humor share: **0.39** · share of comments to top-5 accounts: **0.35** · distinct accounts commented on: **111** · comments that drew a reply: **0.3** · humor (n=32) median views: **61.5** · low_value_repetitive (n=29) median views: **39.5** · generic_agreement (n=28) median views: **61** · promotional (n=28) median views: **63** · original_insight (n=26) median views: **37.5** · relationship_building (n=25) median views: **54**
Algorithm connection	In the published For You pipeline, out-of-network replies are filtered out by default and in-network replies receive the 0.75 out-of-network-style rescore, so a comment is seen mainly inside the parent conversation, not distributed in feeds. Comments can still build mutual-follow relationships, and the reply weight is raised for original posts from mutually-followed authors. The repo has no 'comment boost' for the commenter's own posts. D04 · Out-of-network reposts and replies are filtered out (home-mixer/filters/oon_retweet_reply_filter.rs:15-28) D10 · Out-of-network discount (OonWeightFactor 0.75), also for in-network replies/reposts (xai-value-model/scoring.rs:206-251) B15 · Bidirectional boost eligibility: mutual-follow original posts only (xai-value-model/weights.rs:84-100)
Alternative explanations	Comments within a founder circle may serve partnership/relationship goals rather than reach; Parent-post context (unavailable for most replies) may make short replies more relevant than they look
Impact	Comment exposure is small per reply; quality matters more than volume for building recognition.
Recommendation	Replace generic agreement with 2-3 sentence expert additions on larger in-niche accounts' posts within their first hour; keep relationship comments but stop emoji/one-word replies.
Validation	Track median views and reply-backs per comment for 'expert addition' comments vs the current mix.
Success criteria	Substantive share > 60% and median views per comment above current median.
Confounders	parent-post reach, timing (unmeasured), relationship goals

Weak openings / low-utility structures

statistical association | supported | confidence: medium

Observation	Opening/structure patterns with clear differences: Opening hook type = announcement: 3.71x (n=8); Opening hook type = how_to: 8.87x (n=33); Opening hook type = story: 6.24x (n=8); Narrative structure = promotion: 3.79x (n=10); Narrative structure = tutorial: 6.1x (n=46); Educational value (0-3) = 2.0: 6.97x (n=46); Opening hook type = observation: 0.61x (n=37); Opening hook type = other: 0.22x (n=24); Narrative structure = opinion: 0.35x (n=29); Narrative structure = other: 0.21x (n=28); Educational value (0-3) = 0.0: 0.19x (n=34); Educational value (0-3) = 1.0: 0.45x (n=33).
Evidence	Opening hook type: announcement (n=8) median ratio: **3.71** · Opening hook type: how_to (n=33) median ratio: **8.87** · Opening hook type: observation (n=37) median ratio: **0.61** · Opening hook type: other (n=24) median ratio: **0.22** · Opening hook type: story (n=8) median ratio: **6.24** · Narrative structure: opinion (n=29) median ratio: **0.35** · Narrative structure: other (n=28) median ratio: **0.21** · Narrative structure: promotion (n=10) median ratio: **3.79** · Narrative structure: tutorial (n=46) median ratio: **6.1** · Educational value (0-3): 0.0 (n=34) median ratio: **0.19** · Educational value (0-3): 1.0 (n=33) median ratio: **0.45** · Educational value (0-3): 2.0 (n=46) median ratio: **6.97**
Algorithm connection	Dwell and click-dwell predictions are weighted positively and 'not_dwelled' negatively in published defaults; an opening that earns an expand plausibly raises those predictions. B05 · Published weights: shares (copy link 20), dwell, quote (home-mixer/params/param.rs:381-395) B04 · Published weights: clicks, expands, profile click 0, VQV 0 (home-mixer/params/param.rs:360-380)
Alternative explanations	Hook type correlates with format and topic; Labels are model judgements and can be noisy
Impact	Hooks are cheap to change and testable post-by-post.
Recommendation	Lead with the concrete outcome or how-to promise; avoid opinion-only or vague observational openers.
Validation	Paired test: same topic, two hook styles on different days; compare vs-baseline ratios.
Success criteria	Hook-type gap persists in a controlled 4-week test.
Confounders	topic, format

Reach trend over the window statistical association · mixed · confidence: low

Observation	Median views on authored posts trend at 0.92x per 30 days (90% CI 0.82–1.03, n=157); monthly medians went from 874 (2026-04) to 376 (2026-10).
Evidence	Theil–Sen multiplicative trend per 30d: **0.92** · trend CI low: **0.82** · trend CI high: **1.03** · posts: **157** · first month median views: **874** · last month median views: **376** · implied decline % per 30d: **8**
Algorithm connection	Out-of-network retrieval is driven by viewers' recent engagement histories (Phoenix two-tower, SimClusters seeds). A sustained decline is consistent with fewer viewers having recently engaged with the account, but the repository cannot attribute it. A04 · Viewer's recent engagement sequence feeds Phoenix retrieval (home-mixer/query_hydrators/retrieval_sequence_query_hydrator.rs:36-61) A05 · Phoenix retrieval is a two-tower model over engagement history (phoenix/README.md:99-121) A12 · SimClusters seeds come from the viewer's recent engagements (home-mixer/sources/simclusters_source.rs:222-252)
Alternative explanations	Change in content mix (fewer X Articles or long tutorials in recent months); Seasonality / audience attention (summer vs autumn); Platform-wide changes in view counting or feed composition (weights changed in Sept–Oct; see parameter history); Recent posts' views still accruing despite the 72h maturity filter
Impact	Determines whether the priority is recovering lost distribution or growing from a stable base.
Recommendation	Track weekly median views for each format separately so a format-mix shift is not mistaken for a reach decline; re-run after 4 weeks of monitoring.
Validation	Format-stratified weekly medians from continuous monitoring; check the trend within X Articles only.
Success criteria	Weekly authored median views stop declining for 4 consecutive weeks (format-adjusted).
Confounders	content mix, seasonality, platform weight changes, metric maturity

Several authored posts on the same day statistical association · supported · confidence: low

Observation	Third-or-later authored posts on the same day reach a median 0.23x of baseline (n=9) versus 1.1x for the first post of the day (n=117).
Evidence	1st post of day: median vs-baseline: **1.1** · n first posts: **117** · 3rd+ post of day: median vs-baseline: **0.23** · n 3rd+ posts: **9** · authored posts per week: **6.3** · gap <2h: median vs-baseline (n=5): **0.19** · gap 2-6h: median vs-baseline (n=20): **0.28** · gap 6-24h: median vs-baseline (n=68): **0.6** · gap >24h: median vs-baseline (n=58): **1.7**
Algorithm connection	vm-ranker applies an author-diversity multiplier within each viewer's ranked feed: a repeat author's later candidates are multiplied by (1 - floor) * decay^k + floor with AuthorDiversityDecay 0.5 and AuthorDiversityFloor 0.25. Several fresh posts compete for the same viewer slots. This is per viewer request, so it is a plausible mechanism, not a measured penalty on this account. D09 · Author diversity multiplier decays repeated authors within one feed (xai-value-model/scoring.rs:183-204) B12 · Author diversity decay and out-of-network discount (vm-ranker) (vm-ranker/params.rs:169-228) B13 · ComputeValueModel default false is overridden by the request flag (vm-ranker/scoring/mod.rs:31-35)
Alternative explanations	Later posts of the day may be lower-effort or reactive; Time-of-day differences (later posts land outside peak hours); Small sample for 3rd+ posts
Impact	Spacing posts is a free change; if the association is real, it recovers reach on the weakest third of posts.
Recommendation	Cap authored posts at 2 per day, at least 6 hours apart; move extra content into replies, threads, or the next day.
Validation	2-week alternating schedule (stacked vs spaced days) with the same formats; compare vs-baseline medians.
Success criteria	3rd+-post days disappear without lowering weekly total views.
Confounders	time of day, post quality, format

High-volume templated replies (lead-magnet fulfilment) hypothesis supported confidence: low

Observation	186 template replies ('Sent', 'Here you go: <link>') to 176 distinct accounts, about 7.3 per week. These are fulfilment of comment-for-resource CTAs, inferred from their template text.
Evidence	fulfilment replies: **186** · fulfilment replies per week: **7.3** · distinct recipients: **176** · keyword CTA posts: median views ratio: **1.92** · keyword CTA posts n: **16** · promotional posts: median views ratio: **1.94** · promotional posts n: **23**
Algorithm connection	A published botmaker rule clusters replies with duplicate text and applies a COPYPASTA_SPAM label (skipping high-PageRank and grey-verified users), and the Grox reply_spam flow can label risky replies. Whether these clients' identical 'Sent' replies are labelled, and what any label does to reach, is not observable and partly withheld; this is a risk to monitor, not an observed penalty. C15 · botmaker-rules label duplicate-text replies as COPYPASTA_SPAM (some rules withheld) (botmaker-rules/scarecrow/bot/BBQDuplicateTextRepliesProd.bot:25-46) C11 · Grox reply spam flags replies with RiskyHighVizReply unless high cred (>= 66) (grox/flows/reply_spam/task_write.py:39-80) C12 · Grox prompts withheld; UPA "banger" screen and multimodal embeddings (grox/flows/ptos/prompts.py:11-11)
Alternative explanations	Keyword-CTA posts generate many replies, which is itself a strong positive signal; Templated fulfilment is normal for lead magnets and widely used by peers
Impact	Mainly risk management and time cost; also inflates 'reply' counts so commenting looks larger than it is.
Recommendation	Keep lead magnets but vary fulfilment (short personalised line + resource), or deliver via one pinned reply / DM automation instead of hundreds of identical public replies.
Validation	Compare authored-post reach in weeks with vs without mass fulfilment replies (matched formats).
Success criteria	No drop in lead-magnet conversions with < 25% of current identical public replies.
Confounders	lead-magnet topic quality, CTA wording

External links reduce reach (common claim) statistical association not_supported confidence: low

Observation	Among non-article authored posts, posts containing a t.co URL (external link or media) reach 0.68x the median of those without (n=34 vs 93).
Evidence	non-article posts with t.co URL: ratio: **0.68** · n with URL: **34** · n without URL: **93** · CI: **0.4-1.93**
Algorithm connection	The published ranking code contains no link penalty: open_link carries a small positive weight. has_url is computed only for an experimental debias payload that is off by default. Media and external links cannot be separated in this data. B04 · Published weights: clicks, expands, profile click 0, VQV 0 (home-mixer/params/param.rs:360-380) D17 · Debias payload with has_url is sent to vm-ranker only when enabled (default off) (home-mixer/scorers/vm_ranker_request.rs:17-24)
Alternative explanations	t.co URLs include attached media, which may help; Link posts are often promotional
Impact	Determines whether link placement should change; the repository gives no reason to expect a penalty.
Recommendation	Do not restructure content around a link penalty; if testing, put the link in a self-reply on half of link posts and compare.
Validation	Randomised link-in-post vs link-in-reply on comparable posts for 4 weeks.
Success criteria	Decision rule in experiment plan.
Confounders	media vs link, promotional intent

Signs of visibility limitations insufficient · not_supported · confidence: insufficient evidence

Observation	0 of 157 mature authored posts have fewer than 100 views; reach is continuous with the account's normal range. No public signal indicates a visibility restriction; this audit cannot observe X's visibility-filtering decisions.
Evidence	mature authored posts with < 100 views: **0** · mature authored posts: **157** · median views per reply: **25** · X does not expose visibility labels or filter decisions to the public; none can be observed here
Algorithm connection	Visibility filtering (Allow/Notice/Interstitial/Drop) and account-quality systems exist in the repository, but their per-account outputs are not public. C16 · Visibility-filtering verdicts are Withheld (Drop/Tombstone) or Shown with notices (visibility-filtering/models/verdict.rs:8-26) D13 · VFFilter drops only Drop/Tombstone/NotEvaluated; Downrank/Avoid kept (home-mixer/filters/vf_filter.rs:22-27) C07 · No user-cred input in the For You value-model scoring path (xai-value-model/inputs.rs:3-17)
Alternative explanations	Any decline is better explained by format/topic mix and the trend analysis
Impact	No action warranted on current evidence; claims of suppression would be speculation.
Recommendation	Do not act on shadowban theories; monitor for abrupt step-changes (>50% week-on-week drop across all formats) as the trigger to investigate.
Validation	Monitoring alert on control limits (rolling median ± 3 MAD) across formats.
Success criteria	No unexplained cross-format step change.

8. Algorithm connections

Findings from the published repository referenced in this audit (status: verified = read in code; inferred = reasoned from code). Repository code is not a complete description of live production systems; published parameter defaults are cron-synced but experiments may override them.

ID	Finding	Location	Status
A04	**Viewer's recent engagement sequence feeds Phoenix retrieval** RetrievalSequenceQueryHydrator fetches the viewer's aggregated user-action sequence from the user action aggregation service, with length up to MaxSeqLengthRetrieval (default 1024) and aggregation type PhoenixRetrievalAggregationType (default DENSE_WITH_LONG_DWELL). That sequence is the input PhoenixSource, PhoenixMOESource and PhoenixTopicsSource send to the retrieval model. Relevance: Out-of-network retrieval is driven by what each viewer engaged with, so an account's posts are plausibly more retrievable for viewers whose histories contain similar posts/authors. This is a viewer-side signal, not something an author controls directly. · Limitations: The aggregation service and the exact actions included per aggregation type are not in this repo; the meaning of UAS_WINDOW_TIME_MS passed here is not documented.	home-mixer/query_hydrators/retrieval_sequence_query_hydrator.rs:36-61	verified
A05	**Phoenix retrieval is a two-tower model over engagement history** Per phoenix/README.md, the user tower encodes the viewer's engagement history (plus one user-features token) into a normalized embedding, with no learned per-user ID embedding. The candidate tower embeds posts via semantic IDs plus hashed author IDs; the index is baked into the checkpoint, and serving retrieves top-K by dot product. Relevance: Main ML path by which non-followers can be shown an account's posts: posts whose content/author embeddings sit near a viewer's history embedding are more likely to be retrieved. Content similarity (semantic IDs) and author identity both enter the candidate representation. · Limitations: This is README prose; the code is phoenix/xrex/models/recsys_two_tower_model.py (RecsysTwoTowerModel, RecsysCandidateTower). The README also says retrieval is trained with favorites as the positive signal (phoenix/README.md lines 274-275). Production checkpoints, corpus and cluster configs are not published.	phoenix/README.md:99-121	verified
A07	**Cold-start retrieval pool: low-fav, low-view, fresh posts** cold_pool_mask keeps posts whose metadata shows view_count < max_views (default 500) and fav_count < max_favs (default 8); posts missing metadata are kept. retrieval_dataset.py combines this with fresh_post_mask(cold_start_max_age_seconds) to form the HOME_COLD slice of the home index, and the complement becomes HOME_HOT (DEFAULT_COLD_START_MAX_AGE_SECONDS = 2 * 60 * 60). Relevance: Young posts that have not yet accumulated many views or likes appear to be eligible for a separate cold-start retrieval slice, which could give small accounts' new posts some out-of-network exposure; eligibility appears to end once a post crosses the view/fav/age limits. · Limitations: Defaults are function defaults in training/indexing code (phoenix/xrex); whether production serving passes other values (e.g. cold_start_max_age_seconds) is not shown. HOME_COLD is built from the 1fav_1day snapshot (retrieval_dataset.py line 280), so it presumably requires the post to be in the 1fav index (see A08).	phoenix/xrex/data/cold_pool_filter.py:16-51	verified
A08	**Phoenix retrieval index admission is triggered by favorite events** The rank-all candidate processor skips community posts and reposts, routes replies to separate reply indexes, and skips posts that visibility filtering would drop or that are adult (except an NSFW video index). For other posts, a Fav event source writes the post to the 1fav, 1fav_topic and 32fav (only if favoriteCount >= 32, lines 62-76) indexes; a PostCreation event writes a separate post-creation index. Relevance: Getting at least one like plausibly matters for an original post to enter the 1fav retrieval index that HOME/HOME_COLD slices are built from; posts that VF would drop or that are labelled adult are not indexed there, which would limit non-follower discovery. · Limitations: Which index each Phoenix retrieval cluster actually queries in production is configured outside this file. eventProcessing helpers (shouldDropPostByVF, isNsfwPost) are referenced but their full logic is elsewhere. lifecycle labels differ between strato files (some say Development).	phoenix-rankall-strato/columns/phoenix_rank_all/phoenixRankAllCandidateProcessor.strato:469-530	verified
A11	**SimClusters source constants (score floor, seed age, result caps)** SimclustersSource queries SimClusters ANN with cosine similarity over log-fav embeddings (MODEL_20M_145K_2020), ANN_MAX_NUM_RESULTS = 200 per seed, scanning up to 50 clusters. Results are kept only if score > POST_ANN_MIN_SCORE = 0.5, seeds must be no older than MAX_SEED_AGE (90 days), and the merged output is capped at MAX_RESULTS = 800. Relevance: A second out-of-network path: posts that are cluster-similar (by who favorites them) to posts the viewer engaged with can be surfaced to non-followers. Candidate age is bounded by SimclustersMaxCandidateAgeHours (default 48). · Limitations: The SimClusters embeddings and ANN service are in simclusters/ (Scala) and their training data is not published. The OONNsfwSimclustersFilter later removes SimClusters posts from adult-flagged authors the viewer does not follow.	home-mixer/sources/simclusters_source.rs:25-36	verified
A12	**SimClusters seeds come from the viewer's recent engagements** Seeds are the post ids from the viewer's explicit and implicit engagement signals, sorted by engaged_at_ms (most recent first) and deduplicated; seeds older than MAX_SEED_AGE are dropped. Each seed's ANN results are filtered by POST_ANN_MIN_SCORE, limited to MAX_SANN_CANDIDATES / number_of_seeds per seed, round-robin interleaved, and truncated to 800. Relevance: If viewers who engage with posts similar to an account's posts are numerous, the account's posts may be pulled in as similar candidates; engagement on an account's own posts also becomes a seed for that engager's future SimClusters retrieval. · Limitations: Per-type signal caps (EngagementSignalsMaxPerType, default 15) and the upstream engagement stores (twistly) are outside this file; caching (pod and shared caches, 600s pod TTL) can delay changes.	home-mixer/sources/simclusters_source.rs:222-252	verified
B01	**Predicted heads consumed by the value model** PhoenixScores lists the per-candidate predictions the value model reads: probability heads (favorite, reply, retweet, photo_expand, video_open, click, open_link, profile_click, vqv, share, share_via_dm, share_via_copy_link, dwell, quote, quoted_click, quoted_vqv, follow_author, post_unexplored, react_with_video), negative heads (not_interested, block_author, mute_author, report, not_dwelled) and continuous values (dwell_time, click_dwell_time, home_video_continuation_secs, home_profile_visit_secs). Relevance: These are the only engagement predictions that enter the weighted score; actions not listed (e.g. bookmark) have no direct weight in ranking. · Limitations: The mapping from Phoenix output indices to these fields	xai-value-model/phoenix_scores.rs:1-31	verified

	is in the component_library (PredictNextActions::candidate_scores), not in this repo. Which heads a production checkpoint emits is not published.		
B03	**Published weights: like, reply, bidirectional boost, repost** param.rs sets FavoriteWeight 0.66, ReplyWeight 5.0, BidirectionalFollowReplyWeightBoost 15.0, BidirectionalFollowDwellWeightBoost 0.0 and RetweetWeight 1.0. Each multiplies a predicted probability for the viewer, not an engagement count. Relevance: Posts a viewer is predicted to reply to get relatively large weight per unit probability; the contribution also depends on how likely each action is, which the weights do not show. · Limitations: Line 1 states values are mirrored from config defaults (last sync 2026-10-22T16:00:34Z). README (Experiments and Configuration) says tunable values are read from a configuration system and that cron scripts set code defaults to the primary production values; per-user experiments can still override them.	home-mixer/params/param.rs:345-359	verified
B04	**Published weights: clicks, expands, profile click 0, VQV 0** param.rs sets PhotoExpandWeight 0.05, VideoOpenWeight 0.02, ClickWeight 0.3, OpenLinkWeight 0.2, ProfileClickWeight 0.0 and VqvWeight 0.0. With weight 0.0, predicted profile clicks and video quality views add nothing to the weighted score. Relevance: Under these defaults, predicted profile clicks and video quality views do not raise a post's score; link opens and post clicks contribute small weights. · Limitations: VQV is additionally gated per candidate (vqv_eligible in home-mixer/util/candidates_util.rs: video duration > MinVideoDurationMs and a follower-count condition on the requesting user). README (Experiments and Configuration) says tunable values are read from a configuration system and that cron scripts set code defaults to the primary production values; per-user experiments can still override them.	home-mixer/params/param.rs:360-380	verified
B05	**Published weights: shares (copy link 20), dwell, quote** param.rs sets ShareWeight 2.0, ShareViaDmWeight 5.0, ShareViaCopyLinkWeight 20.0, DwellWeight 0.05 and QuoteWeight 5.0. share_via_copy_link has the largest positive weight in the file. Relevance: Posts viewers are predicted to share (especially via copy link or DM) or quote receive relatively high weight per unit probability; since such actions are presumably rare, the realized effect is uncertain. · Limitations: Weights scale predicted probabilities whose base rates differ by action, so weight size alone does not show relative influence. README (Experiments and Configuration) says tunable values are read from a configuration system and that cron scripts set code defaults to the primary production values; per-user experiments can still override them.	home-mixer/params/param.rs:381-395	verified
B06	**Published weights: follow_author, unexplored, continuous dwell** param.rs sets QuotedClickWeight 0.05, QuotedVqvWeight 0.0, FollowAuthorWeight 4.0, PostUnexploredWeight 0.015, ReactWithVideoWeight 0.0, ContDwellTimeWeight 0.006 and ContClickDwellTimeWeight 0.4. The cont_* weights multiply predicted continuous dwell values rather than probabilities. Relevance: Predicted probability that a viewer follows the author is rewarded, which could favour posts that convert non-followers; predicted dwell also contributes. PostUnexplored applies only to in-network candidates in home-mixer (see B17). · Limitations: Units/scale of predicted dwell_time and click_dwell_time are not documented. README (Experiments and Configuration) says tunable values are read from a configuration system and that cron scripts set code defaults to the primary production values; per-user experiments can still override them.	home-mixer/params/param.rs:396-437	verified
B08	**Weights multiply predicted probabilities; only Home Timeline engagement counts** The param.rs comment says each weight multiplies a predicted probability (or continuous value) and not raw engagement counts, so ratios such as report vs like are not count equivalences. It adds that the baseline report probability is more than 1000x lower than a like, and that for an engagement to count in the recommendation system it must happen on a post served in Home Timeline. Relevance: Coordinated engagement from people navigating directly to a post is described as having no ranking impact; reach depends on per-viewer predicted actions, not totals. · Limitations: The same comment appears in xai-value-model/scoring.rs (lines 28-56) and README.md lines 353-355. The '468' example matches 234/0.5 (the vm-ranker/params.rs favorite default), whereas param.rs now has 0.66 (inferred arithmetic). The training-label pipeline that enforces the Home Timeline rule is not shown in this file.	home-mixer/params/param.rs:316-344	verified
B12	**Author diversity decay and out-of-network discount (vm-ranker)** vm-ranker defaults: EnableAuthorDiversity true with AuthorDiversityDecay 0.5 and AuthorDiversityFloor 0.25, applied as (1 - floor) * decay^k + floor for an author's k-th post in the slate (scoring.rs diversity_multiplier). OonWeightFactor 0.75 multiplies out-of-network posts and, with EnableOonRescoreForInNetworkRepliesRetweets true, in-network replies/reposts; topic requests use TopicOonWeightFactor 0.5. Relevance: A non-follower sees an account's post at a 0.75 multiplier relative to an identical in-network post, and several posts from the same author in one slate get progressively discounted. · Limitations: These params are vm-ranker-only names (not in home-mixer param.rs); home-mixer's local value model disables author diversity and OON rescoring (B17), so on the RPC-failure fallback path these multipliers are not applied. NewUserAgeThresholdSecs defaults to 0, effectively disabling the new-user OON factor. README (Experiments and Configuration) says tunable values are read from a configuration system and that cron scripts set code defaults to the primary production values; per-user experiments can still override them.	vm-ranker/params.rs:169-228	verified

B13	**ComputeValueModel default false is overridden by the request flag** vm-ranker computes the value model when req.compute_value_model is true or the ComputeValueModel param is true (default false in vm-ranker/params.rs). home-mixer's RequestShape::build always sets compute_value_model: true with value_model_id "dpp", so the vm-ranker value model runs regardless of the param default; otherwise vm-ranker would pass through home-mixer's scores. Relevance: Means the vm-ranker weights, author diversity and OON factor (B11/B12) are the ones applied on the normal path, followed by DPP diversity reordering (DppEnabled true, DppTheta 0.65 in vm-ranker/params.rs lines 235-247). · Limitations: If vm-ranker has no resolved config it falls back to upstream scores (reason no_config); home-mixer/scorers/vm_ranker_request.rs lines 37-45 hold the request flag.	vm-ranker/scoring/mod.rs:31-35	verified
B15	**Bidirectional boost eligibility: mutual-follow original posts only** reply_weight_for returns reply + bidirectional_follow_reply_weight_boost when the boost is non-zero and the candidate is bidirectional_boost_eligible, defined in inputs.rs as not a reply, not a repost, and is_mutual_follow_author. Mutual status is hydrated by BidirectionalFollowHydrator (EnableBidirectionalFollowHydration default true). Relevance: With ReplyWeight 5.0 and boost 15.0 the effective reply weight for eligible posts is 20.0 in home-mixer defaults (arithmetic); only original posts benefit, not the author's replies or reposts. · Limitations: Mutual-follow check uses the social graph at request time; failures return errors for the hydrator.	xai-value-model/weights.rs:84-100	verified
C07	**No user-cred input in the For You value-model scoring path** The per-candidate inputs to the value model are predicted engagement scores plus author id, in-network, reply/retweet, mutual-follow and video eligibility flags; there is no user-cred, PageRank or verification field. A search of home-mixer, vm-ranker and xai-value-model finds no reference to user cred or page rank. In bdsm the user_cred_norm feature is filled with zeros. Relevance: In this code, the final weighted score does not appear to apply a direct reputation multiplier. Cred can still matter indirectly through labels whose application it gates (C05/C06/C10/C11), which visibility filtering can act on. · Limitations: This is a negative result from searching the published code. The Phoenix model's own input features are not fully audited here, and a TweetEmbedding thrift struct in phoenix/crates/serving/xai-recsys-mm-server/src/events.rs carries a user_credibility field whose use was not traced.	xai-value-model/inputs.rs:3-17	inferred
C11	**Grox reply spam flags replies with RiskyHighVizReply unless high cred (>= 66)** The reply_spam flow scores replies; when the reply-ranking score is 0.0 it applies the RiskyHighVizReply post label unless the author's user-cred is >= 66 (or the high-PageRank flag when no score), the author has a grey badge, or is a test user. Separately, scarecrow rule 25026 applies RISKY_HIGH_VIZ_REPLY for 14 days when a Grox spamScore >= 0.97, unless the reply is by the conversation author or the author is high-PageRank, skip-listed or gray verified. Relevance: Low-cred accounts that leave replies judged spammy may have those replies labelled; the rule name suggests a downrank of the reply, but where the label is consumed is not in this repo. · Limitations: The consumer of RiskyHighVizReply (e.g. conversation reply ordering) is not present; the effect on reach is inferred from the label name only. Prompts are withheld. See botmaker-rules/scarecrow/bot/GroxTweetProcessor.bot for rule 25026.	grox/flows/reply_spam/task_write.py:39-80	verified
C12	**Grox prompts withheld; UPA "banger" screen and multimodal embeddings** Grox runs LLM/VLM flows on posts: ptos (safety/spam policy), reply_spam, upa (an initial "banger" screen and post-safety screen that send boolean metadata to a labelling endpoint) and mm_emb (multimodal post embeddings). The prompt templates (.j2) are excluded from the repo; each prompts.py says so in a comment. Relevance: Grox outputs can become post labels that visibility filtering acts on, but the classification criteria are not public, so a specific post's risk cannot be predicted from code alone. · Limitations: The README lists Grox prompts among withheld files. How UPA labels are used downstream (grok_upa_action_with_labels) is not shown. mm_emb gates embedding by an adaptive minimum fav count (grox/flows/mm_emb/adaptive_min_fav.py).	grox/flows/ptos/prompts.py:11-11	verified
C15	**botmaker-rules label duplicate-text replies as COPYPASTA_SPAM (some rules withheld)** Rule 21599 consumes scheduled BigQuery jobs that cluster replies with duplicate text (unigrams, or characters for CJK) and applies the COPYPASTA_SPAM label to each reply, skipping test users, high-PageRank users, a temporary skip list and gray-verified users. Some botmaker rules are deliberately not published. Relevance: Posting the same or near-identical reply text repeatedly could plausibly get replies labelled; the rule's own note calls this a downrank, but the label's ranking effect is not shown in this repo. · Limitations: README.md line 433 lists "Some botmaker rules" as withheld. The clustering query and its similarity thresholds are not published.	botmaker-rules/scarecrow/bot/BBQDuplicateTextRepliesProd.bot:25-46	verified
C16	**Visibility-filtering verdicts are Withheld (Drop/Tombstone) or Shown with notices** The visibility-filtering service returns either Withheld (Drop or Tombstone) or Shown with optional notice (soft intervention, appealable notice, label interstitial), media restriction and limited-engagement parts. Per the README, the first rule that answers drop ends evaluation. Relevance: Only drop-type outcomes remove a post from For You in home-mixer (see D13); notices and interstitials change presentation, not eligibility, in the code read here. · Limitations: The mapping from Verdict to the client Action enum used by home-mixer (Allow, Interstitial, Avoid, Downrank, Drop, Tombstone, NotEvaluated) was not fully traced.	visibility-filtering/models/verdict.rs:8-26	verified
D04	**Out-of-network reposts and replies are filtered out** Reposts from accounts the viewer does not follow are removed, as are out-of-network replies unless EnablePhoenixOonReplies is on and the reply came from Phoenix retrieval (optionally only replies to the root). Any reply whose ancestors failed to hydrate is removed regardless of network. Relevance: Replies and reposts mostly reach For You only for the author's followers; with EnablePhoenixOonReplies defaulting to false, an account's replies are unlikely to be recommended to non-followers through this pipeline. · Limitations: EnablePhoenixOonReplies default (false) is from home-mixer/params/param.rs; experiments can enable it for some viewers.	home-mixer/filters/oon_retweet_reply_filter.rs:15-28	verified

D09	**Author diversity multiplier decays repeated authors within one feed** Candidates are ordered by score; the k-th additional post from the same author in that request (k = 0 for the first) is multiplied by (1 - floor) * decay^k + floor. With defaults decay 0.5 and floor 0.25 this gives 1.0, 0.625, 0.4375, ... approaching 0.25. Relevance: Several posts by the same author competing for one viewer's feed are progressively discounted; this is per viewer per request, not a global penalty on posting frequency. · Limitations: Defaults are in vm-ranker/params.rs (lines 169-186). Home-mixer's local fallback scoring disables author diversity (home-mixer/scorers/value_model.rs).	xai-value-model/scoring.rs:183-204	verified
D10	**Out-of-network discount (OonWeightFactor 0.75), also for in-network replies/reposts** The OON multiplier applies to out-of-network candidates and, when EnableOonRescoreForInNetworkRepliesRetweets is true (default), to in-network replies and reposts. The multiplier is the request's effective_oon_weight: OonWeightFactor (0.75) normally, TopicOonWeightFactor (0.5) for topic requests, or NewUserOonWeightFactor for eligible new viewers. Relevance: Reaching non-followers carries a 25% score discount at defaults, so out-of-network posts need higher predicted engagement to rank equally. · Limitations: Defaults from vm-ranker/params.rs lines 205-228; production values can differ per experiment bucket.	xai-value-model/scoring.rs:206-251	verified
D13	**VFFilter drops only Drop/Tombstone/NotEvaluated; Downrank/Avoid kept** After selection, VFFilter removes candidates whose visibility action is Drop, Tombstone or NotEvaluated. Allow, Interstitial, Avoid and Downrank are kept, and no other home-mixer code was found that changes scores based on Avoid or Downrank. Relevance: Visibility filtering affects For You reach mainly by dropping posts. A post whose VF lookup failed (NotEvaluated) would also be dropped. · Limitations: The absence of Downrank/Avoid handling is from a search of home-mixer only; other services could use them.	home-mixer/filters/vf_filter.rs:22-27	verified
D17	**Debias payload with has_url is sent to vm-ranker only when enabled (default off)** When VMRankerSendDebiasInputs is true and the request has no cached posts, home-mixer adds per-candidate content features to the vm-ranker request: has_video, video duration, has_photo, media count, weighted text length, newline count and has_url (external URLs in text). The parameter defaults to false, and the value model shown here does not read these fields. Relevance: No URL penalty appears in the published scoring formula; has_url is only an experimental payload field whose use by vm-ranker is not shown. · Limitations: Default from home-mixer/params/param.rs lines 502-507. How vm-ranker experiments consume experiment_payload is not in this repo.	home-mixer/scorers/vm_ranker_request.rs:17-24	verified
D19	**Cold-start lift for low-impression posts from smaller authors** AuthorColdStart can lift one eligible candidate per request to the score at a randomly chosen target rank. Eligible posts are originals (not replies or reposts) from authors with <= 50,000 followers, younger than 7,200 s, with fewer than 200 home impressions, within an allowed position range, and in the viewer's experiment arm's author corpus. Relevance: Small accounts' fresh original posts can get an exploratory boost early in their life; replies and reposts are not eligible. · Limitations: Defaults from home-mixer/params/param.rs (lines 539-581); slot range and corpus bucketing are experiment-controlled.	home-mixer/scorers/author_cold_start.rs:188-193	verified

9. Competitive comparison

Same slice for client and peers: authored posts, last 90 days, mature only.

Account	Followers	Posts/wk	Median views	Views / 1k followers	Eng /1k views	Replies /1k	Bookmarks /1k	Article share
@nifinet	**11,613**	**7.6**	**1,183**	**101.9**	**9.2**	**0.93**	**2.14**	**14%**
@jspujji	181,489	2.7	6,251	34.4	7.4	0.91	0.83	0%
@lorenzo_pravata	27,775	14.3	3,237	116.5	24.7	1.93	10.96	38%
@jordan_ross_8F	90,498	15.9	2,917	32.2	7.5	1.33	0.91	2%
@georgeclem	19,035	17.1	929	48.8	25.5	3.18	5.30	9%
@wickedguro	25,570	11.5	8,009	313.2	20.8	3.50	2.41	2%
@RodmanAi	35,125	13.4	6,698	190.7	41.2	1.73	18.35	0%

- Peers: newest <=100 original posts per peer in the window were collected; posting rate is computed over the span those posts cover.
- Size normalisation: views per 1,000 followers uses follower counts observed on the snapshot date, applied to the whole window (no follower history).
- Engagement per 1,000 views is size-independent but depends on audience composition.
- Peer reply/comment behaviour was not collected; the commenting audit has no peer benchmark.
- Quote-post share is not compared: peers were collected with the console's 'original' filter.
- Peer topic mix uses the heuristic labeller (keywords), not the reviewed labels used for clients.

10. Post rewrites (editorial proposals to test)

> A rewrite is an editorial proposal; it is not predicted to receive more distribution. Each comes with a measurable test.

1. quote post · 2026-10-10 17:25 UTC · 181 views (0.07x baseline)

> Original: the scariest competitor for a saas in 2026 is your own customer opening claude code on a saturday.
>
> if your price isn't tied to the value, someone will vibe code the gap.

Observed issue	Reached 0.07x of the trailing baseline (181 views vs a 2670 median for the account's previous 28 days). It drew 0 like(s), 1 reply, 0 reposts and 0 bookmark(s) (5.5 engagements and 5.5 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (5.5/1k) was below the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'neither'.
Cannot conclude	Why it underperformed. Timing (posted Saturday 17:25 UTC, post 2 of 3 from this sample that day; other posts that day are not in the input), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions.

> The scariest competitor for a SaaS in 2026 is your own customer opening Claude Code on a Saturday.
>
> If your price isn't tied to value, someone will vibe code the gap.
>
> Check your pricing:
> 1. Does it move when the customer's outcome moves?
> 2. Could they rebuild it in a weekend?

Editorial changes	• Turned the quote post into a standalone post so it does not depend on the quoted tweet for context • Kept the original hook and both of Nicolas's claims word for word in substance • Added a two-question self-check so the reader leaves with a concrete step, matching the audit's finding that how-to posts do better on this account • Capitalised the opening for readability; no new numbers or claims added
Test	Publish as a standalone post (not a quote post) as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.07x, 5.5 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.
Drafted by	claude-sonnet (session reviewer)

2. long post · 2026-07-08 22:30 UTC · 202 views (0.07x baseline)

> Original: Most AI outbound is just bad outbound with a bigger engine.
>
> If your sequence did not book meetings when you sent it manually, automating it will not change that.
>
> You just send the same thing to more people, faster.
>
> Automating it cannot add judgment you do not have. It will not notice that your offer is weak, or that you are reaching out three months too late. Those are still your calls to get right, and the agent assumes you already did.
>
> So get the sequence working by hand first. Find the angle and the timing that convert, and automate it once it does.
>
> Whatever you hand the agent is what it scales, so hand it something that already works.

Observed issue	Reached 0.07x of the trailing baseline (202 views vs a 2942 median for the account's previous 28 days). It drew 4 like(s), 0 replies, 0 reposts and 1 bookmark(s) (24.8 engagements and 0.0 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (24.8/1k) was above the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'resonance without reach'.
Cannot conclude	Why it underperformed. Timing (posted Wednesday 22:30 UTC), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions.

> Automating a sequence that never booked meetings by hand just sends the same thing to more people, faster.
>
> The agent can't add judgment you don't have. It won't notice that your offer is weak, or that you're reaching out three months too late. It assumes you already got those calls right.
>
> So here's the order I'd run it in:
>
> 1. Send the sequence by hand first.
> 2. Find the angle that gets replies.
> 3. Find the timing that converts.
> 4. Only then hand it to the agent.
>
> Whatever you hand the agent is what it scales. Hand it something that already works.

Editorial changes	• Replaced the opinion-only opener ('Most AI outbound is just bad outbound...') with the concrete consequence, since opinion-led openers underperform on this account • Converted the closing advice into four numbered steps so the post reads as a how-to (the format the audit found strongest here) • Kept every original fact and phrase that carries the argument (weak offer, three months too late, 'whatever you hand the agent is what it scales') • Shortened slightly and removed repetition; no numbers or results added
Test	Publish as a long post as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.07x, 0.0 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.
Drafted by	claude-sonnet (session reviewer)

3. long post · 2026-07-06 17:00 UTC · 239 views (0.08x baseline)

Original: Bought lead lists are the most overrated asset in B2B. The data is fine the day you buy it. It goes stale the moment after, and you keep working it like it is still live.

Think about what a list actually is.

A name, a title, a company, and a timestamp from whenever someone scraped it. It tells you who a person was on the day they got added, and nothing about what is happening in their market today.

And this week is what matters.

Someone who looked perfect six months ago might have just signed with your competitor.

Someone who ghosted you last quarter might have raised a round on Tuesday and is buying everything in sight.

The list knows none of it. It cannot. It is a photo of a moment that already passed.

Watch how the buy side of a market behaves and you notice what a list cannot show.

Demand never sits still.

Companies start looking for a vendor every day, well before they land on anyone's radar. That motion is the only real signal, and a bought list has none of it.

So with a list you do the only thing it lets you do. You blast everyone and hope. You send 1,000 emails to reach the 20 people who happen to be ready, and you burn the other 980 to get there.

Signal-based outbound runs the other way.

You start from the motion.

Someone opens a role they only open right before they buy. Someone drops the tool you replace. Someone closes a round.

That is the moment a stranger becomes worth a message.

Then the real work starts, and the work is restraint more than anything else. Score the why-now, pick the handful that matter, let the rest be. Only then does someone get an email, written off the real trigger.

TL;DR: A bought lead list just tells you who people were. The market tells you who is ready now. Build from the second one.

Observed issue	Reached 0.08x of the trailing baseline (239 views vs a 2956.5 median for the account's previous 28 days). It drew 4 like(s), 0 replies, 0 reposts and 1 bookmark(s) (20.9 engagements and 0.0 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (20.9/1k) was above the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'resonance without reach'.
Cannot conclude	Why it underperformed. Timing (posted Monday 17:00 UTC), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions.

A bought lead list makes you send 1,000 emails to reach the 20 people who happen to be ready. You burn the other 980 to get there.

The data is fine the day you buy it. It goes stale right after, and you keep working it like it's still live. A list is a name, a title, a company, and a timestamp from whenever someone scraped it. It tells you who a person was, and nothing about what's happening in their market this week.

And this week is what matters. Someone who looked perfect six months ago might have just signed with your competitor. Someone who ghosted you last quarter might have raised a round on Tuesday and is buying everything in sight.

Here's how to build from the motion instead of the list:

1. Pick the moments that come right before a purchase. Someone opens a role they only open right before they buy. Someone drops the tool you replace. Someone closes a round.

2. Watch for those moments. Demand never sits still, and companies start looking for a vendor well before they land on anyone's radar.

3. Score the why-now on every account that moves.

4. Pick the handful that matter and let the rest be. This is where most of the work is: restraint.

5. Only then send an email, written off the real trigger.

TL;DR: A bought lead list tells you who people were. The market tells you who is ready now. Build from the second one.

Editorial changes	• Moved the most concrete line (1,000 emails to reach 20, burning 980) from mid-body to the opening • Restructured the signal-based half into five numbered steps, turning an opinion essay into a how-to • Kept the original examples (role opened, tool dropped, round closed, signed with competitor, raised on Tuesday) and the TL;DR • Cut the 'most overrated asset' framing, an opinion-led opener of the type that underperforms on this account • Kept similar long-form length; no new numbers, clients or results
Test	Publish as a long post as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.08x, 0.0 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.
Drafted by	claude-sonnet (session reviewer)

4. quote post · 2026-10-10 09:57 UTC · 330 views (0.12x baseline)

> Original: Imagine tu as le choix entre:
>
> - try hard sur vendre et build du produit pour faire du MRR
>
> - try hard d’économiser 99€/mois (3€/jours)
>
> quel est le meilleur ROI?

Observed issue	Reached 0.12x of the trailing baseline (330 views vs a 2671.5 median for the account's previous 28 days). It drew 1 like(s), 1 reply, 0 reposts and 0 bookmark(s) (6.1 engagements and 3.0 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (6.1/1k) was below the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'neither'.
Cannot conclude	Why it underperformed. Timing (posted Saturday 09:57 UTC, post 1 of 3 from this sample that day; other posts that day are not in the input), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions.

> Tu passes combien de temps à économiser 99€/mois (3€/jour)?
>
> Ce temps ne va ni à la vente ni au produit, alors que c'est ça qui fait le MRR.
>
> L'économie est plafonnée à 99€/mois. Le MRR, non.
>
> Ton dernier abonnement sur lequel tu as perdu une heure, c'était lequel?

Editorial changes	• Kept the post in French, as in the original • Made it standalone instead of a quote post, so the trade-off is readable without the quoted tweet • Stated the implied answer (selling and building beat cutting a 99€ cost) instead of leaving a bare poll-style question; Nicolas should confirm this is the answer he meant • Ended on a specific, easy-to-answer question to invite replies, since replies per 1k are below peers; no generic 'what do you think' closer • Kept only the original numbers (99€/mois, 3€/jour)
Test	Publish as a standalone post (not a quote post) as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.12x, 3.0 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.
Drafted by	claude-sonnet (session reviewer)

5. quote post · 2026-10-10 19:47 UTC · 278 views (0.14x baseline)

> Original: the next buying signal isn't a human on your pricing page.
> it's their agent comparing you to 4 competitors at 2am.
> whoever learns to read agent visits first eats everyone's outbound lunch.

Observed issue	Reached 0.14x of the trailing baseline (278 views vs a 2059 median for the account's previous 28 days). It drew 0 like(s), 1 reply, 0 reposts and 0 bookmark(s) (3.6 engagements and 3.6 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (3.6/1k) was below the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'neither'.
Cannot conclude	Why it underperformed. Timing (posted Saturday 19:47 UTC, post 3 of 3 from this sample that day; other posts that day are not in the input), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions. The claim that agent visits are a readable signal is Nicolas's opinion, not something this data tests.

> A buying signal most teams aren't reading yet: a buyer's agent comparing you to 4 competitors at 2am.
>
> Whoever learns to read agent visits first eats everyone's outbound lunch.
>
> Where I'd start: check which of your pages get read by agents, and when.

Editorial changes	• Removed the 'isn't a human... it's their agent' construction, which the voice rules ban as an AI tell • Opened with the concrete scene (agent, 4 competitors, 2am) instead of the contrarian setup • Added one concrete first step so the post ends with an action, in line with the how-to finding • Standalone instead of a quote post; no new numbers or capability claims (it does not claim any tool detects agent visits)
Test	Publish as a standalone post (not a quote post) as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.14x, 3.6 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.
Drafted by	claude-sonnet (session reviewer)

6. thread · 2026-06-24 20:43 UTC · 315 views (0.14x baseline)

> Original: Not every buying signal deserves a human.
>
> Most teams treat every signal the same. A content download and a week of pricing-page visits land in the same queue, a rep works them in whatever order, and the hot accounts go cold while someone writes a careful note to a tire-kicker.
>
> The solution is a routing rule. Score each account by readiness, then send it to the right responder.
>
> Human, today: the pricing page hit repeatedly this week, a competitor or review-site comparison, or an inbound reply. High intent and time-sensitive. A person sends a short, specific note.
>
> Agent: a role opened that maps to a pain you solve, or a tech or integration change. Real but early. An agent runs a value-first touch, no pitch.
>
> Nurture: a single download or a newsletter open. Curiosity, not intent. It goes on an educational drip until something stronger shows up.
>
> Ignore: no fit, or every signal is stale. Suppress it so it stops eating attention.
>
> A couple of rules keep it honest. When signals conflict, route to the higher tier and say why. And never act on a signal that isn't actually there.
>
> That's the whole rubric. Here's a simple prompt to help you get started (check the next tweet):

Observed issue	Reached 0.14x of the trailing baseline (315 views vs a 2261 median for the account's previous 28 days). It drew 2 like(s), 3 replies, 0 reposts and 1 bookmark(s) (19.0 engagements and 9.5 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (19.0/1k) was above the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'resonance without reach'.
Cannot conclude	Why it underperformed. Timing (posted Wednesday 20:43 UTC), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions. It was a thread; whether the 'check the next tweet' hand-off affected views of the first tweet is not observable from these fields.

> Not every buying signal deserves a human. Here's a routing rule to decide who responds to each one.
>
> Score each account by readiness, then send it to one of four responders:
>
> 1. Human, today. The pricing page hit repeatedly this week, a competitor or review-site comparison, or an inbound reply. High intent and time-sensitive. A person sends a short, specific note.
>
> 2. Agent. A role opened that maps to a pain you solve, or a tech or integration change. Real but early. An agent runs a value-first touch, no pitch.
>
> 3. Nurture. A single download or a newsletter open. Curiosity, without intent yet. It goes on an educational drip until something stronger shows up.
>
> 4. Ignore. No fit, or every signal is stale. Suppress it so it stops eating attention.
>
> Two rules keep it honest. When signals conflict, route to the higher tier and say why. And never act on a signal that isn't actually there.
>
> Without a rule, a content download and a week of pricing-page visits land in the same queue, and the hot accounts go cold while someone writes a careful note to a tire-kicker.
>
> The prompt to run this rubric is in the reply below.

Editorial changes	• Promised the routing rule in the opening so the reader knows a usable framework follows • Numbered the four tiers so the post scans as a step-by-step rubric • Moved the 'everything in one queue' problem after the rubric so the useful part comes first • Kept the original tiers, examples and the two honesty rules unchanged in substance • Kept the prompt as a reply but made the main post self-contained, so it reads fine even if the reply is not opened; the prompt text itself is not in the input and must come from the original thread
Test	Publish as a long post as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.14x, 9.5 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.

Drafted by	claude-sonnet (session reviewer)

7. quote post · 2026-05-05 08:06 UTC · 157 views (0.14x baseline)

> Original: Package this into a prompt or a .md, it’s not even the point.
>
> @pmarca post if more about the substance, uncovering blind spots, and guiding the LLM into the best outcomes. Give it to any AI and they’ll figure it out

Observed issue	Reached 0.14x of the trailing baseline (157 views vs a 1123 median for the account's previous 28 days). It drew 0 like(s), 0 replies, 0 reposts and 0 bookmark(s) (0.0 engagements and 0.0 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views and drew no engagement at all (0.0/1k), the lowest in this rewrite set. Quadrant label: 'neither'.
Cannot conclude	Why it underperformed. Timing (posted Tuesday 08:06 UTC), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions. The original relied on the quoted post for context; how much that limited reach is not observable here.

> Prompt or .md file, the packaging is not the point.
>
> The substance: uncover blind spots, guide the LLM to the best outcome.
>
> Before formatting, write down what you want, what you're likely missing, and what a great answer looks like. Any AI will figure out the rest.

Editorial changes	• Made it standalone so the insight does not depend on the quoted @pmarca post (if kept as a quote post, fix the 'post if more' typo to 'post is more') • Led with the claim instead of the observation-style reaction, a format that underperforms on this account • Turned 'blind spots' and 'best outcomes' into a short, concrete pre-prompt checklist • No numbers, names or results added
Test	Publish as a standalone post (not a quote post) as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.14x, 0.0 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.
Drafted by	claude-sonnet (session reviewer)

8. long post · 2026-08-14 18:33 UTC · 422 views (0.15x baseline)

> Original: Two signals at the same account beat any single signal. Only when they are different kinds, though.
>
> Say someone at a target account starts engaging your competitor's posts. On its own, mild. Now a role opens on that team to own the exact problem you solve. Also mild.
>
> Together they stop being mild. One says they are shopping. The other says somebody is now accountable for fixing it.
>
> Two different questions, both answered in the same month.
>
> Compare that to two social signals. One person likes a post, someone else comments a week later. It feels like momentum, and it is the same fact twice.
>
> So how do you know which signals to act on and which ones to ignore? Simple heuristic: Do your two signals answer different questions, or the same question louder?
>
> Most setups reward the second kind, because social fires constantly and a pair of them is easy to find.
>
> Building the instinct for your own market doesn't take much time, honestly.
>
> > Just take your last ten closed deals.
> > Go back thirty days before the first real conversation and write down everything that was happening at that account.
>
> You are sure to find patterns and signals that may appear in pairs. Make a note of those, and maybe try watching out for such signals that come in combinations.

Observed issue	Reached 0.15x of the trailing baseline (422 views vs a 2874.5 median for the account's previous 28 days). It drew 4 like(s), 1 reply, 0 reposts and 3 bookmark(s) (19.0 engagements and 2.4 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (19.0/1k) was above the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'resonance without reach'.
Cannot conclude	Why it underperformed. Timing (posted Friday 18:33 UTC), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions.

> How to find the signal pairs worth acting on, using your own closed deals:
>
> 1. Take your last ten closed deals.
> 2. Go back thirty days before the first real conversation at each account.
> 3. Write down everything that was happening there.
> 4. Mark the signals that showed up together.
> 5. For each pair, ask: do these two answer different questions, or the same question louder?
>
> Step 5 is the one that matters. Two signals at the same account beat any single signal, but only when they are different kinds.
>
> Say someone at a target account starts engaging your competitor's posts. On its own, mild. Then a role opens on that team to own the exact problem you solve. Also mild. Together, one says they are shopping and the other says somebody is now accountable for fixing it. Two different questions, both answered in the same month.
>
> Two social signals work differently. One person likes a post, someone else comments a week later. It feels like momentum, and it is the same fact twice. Most setups reward this kind, because social fires constantly and a pair is easy to find.
>
> Keep the pairs that answer different questions, and watch for those combinations going forward.

Editorial changes	• Opened with the method (the audit's best-performing format here) instead of an observation-style opener • Moved the closed-deals exercise from the end to the top and split it into five numbered steps • Kept the competitor-engagement + new-role example and the 'same fact twice' contrast as the explanation for step 5 • Replaced the hedged ending ('maybe try watching out') with one direct instruction • Kept similar long-form length; only the original quantities (ten deals, thirty days, a week, same month) are used
Test	Publish as a long post as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.15x, 2.4 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.

Drafted by	claude-sonnet (session reviewer)

9. long post · 2026-06-02 20:59 UTC · 193 views (0.15x baseline)

Original: If you're running cold outbound on AI tools right now and your reply rates are flat, you're probably building the stack in the wrong order.

Here's the order that actually works:

1. Signals first.

> The triggers worth watching: reposted jobs, fundraising rounds, M&A filings, tenders and RFPs, competitor engagement.
> Each one marks a different buying moment.
> Pick four and wire them in before you touch anything else.

2. Scoring next.

> ICP fit x signal recency x buying-stage weight.
> A reposted job in the last seven days at an ICP-fit account beats a job posting from six weeks ago at a non-ICP account.
> Decide who to act on this week, then who falls into next month.

3. Message context last.

> Pass the exact role being reposted, the exact funding amount, the exact tender description into whatever writes the first line.
> Generic "fits your ICP" filler tells the model nothing useful.

Each layer of this stack feeds the next.

Signal triggers scoring > scoring filters who matters this week > context shapes the message itself

Run this loop.

Observed issue	Reached 0.15x of the trailing baseline (193 views vs a 1302 median for the account's previous 28 days). It drew 2 like(s), 0 replies, 0 reposts and 2 bookmark(s) (20.7 engagements and 0.0 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (20.7/1k) was above the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'resonance without reach'.
Cannot conclude	Why it underperformed. Timing (posted Tuesday 20:59 UTC), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions. This post was already a how-to with steps, so format alone does not explain the gap here.

Build your AI outbound stack in this order. If reply rates are flat, check whether you started at the wrong end.

1. Signals first.
Pick four triggers and wire them in before you touch anything else. Options: reposted jobs, fundraising rounds, M&A filings, tenders and RFPs, competitor engagement. Each one marks a different buying moment.

2. Scoring next.
ICP fit x signal recency x buying-stage weight. A job reposted in the last seven days at an ICP-fit account beats a job posting from six weeks ago at a non-ICP account. Decide who you act on this week, and who falls into next month.

3. Message context last.
Pass the exact role being reposted, the exact funding amount, the exact tender description into whatever writes the first line. Generic "fits your ICP" filler tells the model nothing useful.

Each layer feeds the next: the signal triggers scoring, scoring filters who matters this week, context shapes the message.

Run this loop.

Editorial changes	• Opening now states the instruction first; the conditional 'if you're running... you're probably...' setup moved to a short second sentence • Put the action ('pick four and wire them in') at the start of step 1 instead of the end • Kept the three-step structure, the scoring formula and the seven-days vs six-weeks example unchanged in substance • Tightened the closing loop line; no new numbers or claims
Test	Publish as a long post as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.15x, 0.0 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.

Drafted by	claude-sonnet (session reviewer)

10. quote post · 2026-08-15 16:43 UTC · 443 views (0.15x baseline)

> Original: good reminder
>
> it doesn’t matter where you start
>
> what matters is that you keep showing up and shipping for the customers

Observed issue	Reached 0.15x of the trailing baseline (443 views vs a 2873 median for the account's previous 28 days). It drew 0 like(s), 0 replies, 0 reposts and 0 bookmark(s) (0.0 engagements and 0.0 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views and drew no engagement at all (0.0/1k), the lowest in this rewrite set. Quadrant label: 'neither'.
Cannot conclude	Why it underperformed. Timing (posted Saturday 16:43 UTC), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions. With zero engagements, there is no signal at all about how readers received the content.

> It doesn't matter where you start.
>
> What matters is that you keep showing up and shipping for the customers.
>
> What did you ship for a customer this week?

Editorial changes	• Standalone instead of a quote post, so it reads without the quoted tweet • Dropped the 'good reminder' preamble so the first line carries the idea • Added a specific question to give readers an easy reason to reply • Editor note: this stays an opinion-only short post, the format that underperforms here; it is stronger if Nicolas adds one real thing Sortlist shipped this week (not invented in this draft)
Test	Publish as a standalone post (not a quote post) as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.15x, 0.0 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.
Drafted by	claude-sonnet (session reviewer)

11. quote post · 2026-10-12 15:46 UTC · 209 views (0.18x baseline)

> Original: Just did a cold call session on the warm intent list
>
> 80% connection rate (I dial twice every time).
>
> 15% meeting booked 😳
>
> this is stroooooong

Observed issue	Reached 0.18x of the trailing baseline (209 views vs a 1133.5 median for the account's previous 28 days). It drew 3 like(s), 0 replies, 0 reposts and 1 bookmark(s) (19.1 engagements and 0.0 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (19.1/1k) was above the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'resonance without reach'.
Cannot conclude	Why it underperformed. Timing (posted Monday 15:46 UTC), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions. The original does not state how many calls were made, so the 80% and 15% figures cannot be judged for reliability.

> Just did a cold call session on the warm intent list.
>
> 80% connection rate
> 15% meeting booked 😳
>
> What I did:
> 1. Called only accounts on the warm intent list
> 2. Dialed twice every time
>
> This is strong. Do you call intent accounts before or after the first email?

Editorial changes	• Kept both original results (80% connection, 15% meeting booked) and put them on their own lines up top • Turned the method that was buried in parentheses ('I dial twice every time') into explicit steps, so the post shows how, not just what • Ended with a specific question to invite replies, since replies per 1k are below peers • Standalone instead of a quote post; kept the single emoji; no sample size or other numbers added (the original gives none, so the rates should be read as one session)
Test	Publish as a standalone post (not a quote post) as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.18x, 0.0 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.
Drafted by	claude-sonnet (session reviewer)

12. quote post · 2026-10-11 07:00 UTC · 252 views (0.19x baseline)

> Original: been running a marketplace for 12 years that's literally this
>
> humans pick the prettiest profile on page 1
> agents gonna read all 400 case studies at 2am and pick the one that actually fits
>
> kinda terrifying. also kinda fair ngl

Observed issue	Reached 0.19x of the trailing baseline (252 views vs a 1358.5 median for the account's previous 28 days). It drew 0 like(s), 0 replies, 0 reposts and 1 bookmark(s) (4.0 engagements and 0.0 replies per 1k views).
Can conclude	It underperformed this account's own baseline on views. Its engagement rate (4.0/1k) was below the median of the 12 posts in this rewrite set (12.5/1k); the input does not include the full-account engagement median. Quadrant label: 'neither'.
Cannot conclude	Why it underperformed. Timing (posted Sunday 07:00 UTC), audience mix, competing posts and ranking internals are unobserved, and one post is not evidence of a pattern. With view counts this small, the engagement rate rests on very few actions.

> Been running a marketplace for 12 years. It's literally this.
>
> Humans pick the prettiest profile on page 1.
> Agents read all 400 case studies at 2am and pick the one that fits.
>
> Kinda terrifying. Also fair.
>
> Selling on a marketplace? Write for the reader who reads everything.

Editorial changes	• Standalone instead of a quote post; first line keeps his unique vantage (12 years running a marketplace) • Kept the human-vs-agent contrast and his 'terrifying / fair' reaction • Added one concrete takeaway for agencies and vendors so the post ends with an action instead of an observation • Dropped 'ngl' and 'gonna' for readability; no new numbers or claims
Test	Publish as a standalone post (not a quote post) as the first post of the day on a weekday between 12:00 and 17:00 UTC, with no other post within the prior few hours. After 7 days compare views/baseline and replies per 1k against the original (0.19x, 0.0 replies/1k), plus bookmarks per 1k. Text, format and timing all change at once, so a single rerun cannot attribute any difference to the copy alone; treat it as one data point.
Drafted by	claude-sonnet (session reviewer)

---

11. 30-day growth plan

> Sample sizes assume a two-sided test at alpha 0.05 and 80% power on log views, using the residual spread after month/format/topic controls. Effects smaller than ~1.3x cannot be detected within 30 days at current posting rates. Needs ~393 posts per arm; at the current 6.3 authored posts/week that is ~126 weeks.

Week	Experiment	Hypothesis · intervention · control	Primary metric	Decision rule	Power
1	**Freeze the baseline and fix data gaps**	Current medians are a valid benchmark for the next 3 weeks. *Do:* Start daily metric snapshots (monitoring) for posts and follower count; log every post's format and topic at publish time. *Compare:* n/a Secondary: metric snapshots per post at 24h/72h/7d · Confounders: none · Minimum period: 7 days	data completeness	Proceed when 24h/72h snapshots exist for every new post.	n/a
1	**Space authored posts (max 2/day, >= 6h apart)** diagnostic: supported, low	Same-day stacking is associated with lower reach for later posts (author-diversity mechanism). *Do:* Cap authored posts at 2 per day, >= 6h apart; extra items become replies/threads or move to the next day. *Compare:* Previous 4 weeks of the same account (pre/post) and the stacking days in history Secondary: weekly total authored views, posts/week · Confounders: format mix, seasonality · Minimum period: 3-4 weeks	median views relative to trailing baseline	Keep spacing if median vs-baseline of posts that would have been 3rd+ rises above 0.8x without lowering weekly total views by > 10%.	165 posts/arm **underpowered in 30 days (~52.7 wks)** — read as directional
2	**Article-anchored week vs mixed week** diagnostic: supported, high	X Articles (with a native preview post) outperform other formats on matched topics. *Do:* Alternate weeks: anchor week = 2-3 X Articles on core topics + supporting posts; control week = usual mix. *Compare:* Alternating control weeks (A/B/A/B) Secondary: bookmarks per 1k views, replies per 1k views · Confounders: topic, promotion effort, co-founder reshares · Minimum period: 4 weeks (alternating)	median views per authored post	Adopt as default if anchor-week median >= 1.3x control-week median in both alternating pairs.	393 posts/arm **underpowered in 30 days (~125.5 wks)** — read as directional
2	**Outcome-first hooks vs current openers** diagnostic: supported, medium	Openings that state a concrete outcome or how-to promise earn more expands/dwell than observational openers. *Do:* Randomise (coin flip per post) between an outcome/how-to hook and the author's natural opener for non-article posts. *Compare:* Coin-flip control posts Secondary: likes per 1k, bookmarks per 1k · Confounders: topic, format · Minimum period: 3-4 weeks	median views vs baseline	Adopt if treatment median vs-baseline exceeds control by >= 1.3x with both arms >= 8 posts.	165 posts/arm **underpowered in 30 days (~52.7 wks)** — read as directional
2	**Specific closing prompt + first-hour author replies** diagnostic: supported, medium	Ending with a specific, answerable question and replying to early commenters raises replies per 1k views. *Do:* Half of posts (randomised) end with a specific question; author replies to the first 5 comments within 60 minutes. *Compare:* Randomised posts without the prompt Secondary: views vs baseline, share of posts with >= 5 replies · Confounders: topic, time of day · Minimum period: 4 weeks	replies per 1k views	Adopt if replies per 1k views is >= 1.3x control with no drop in views vs baseline.	165 posts/arm **underpowered in 30 days (~52.7 wks)** — read as directional
3	**Expert-addition comments on larger in-niche accounts** diagnostic: supported, medium	2-3 sentence expert additions on in-niche posts earn more views and reply-backs than generic replies. *Do:* 10 comments/day: substantive additions on posts < 60 min old from accounts with 2-20x the client's	median views per comment	Continue if median views per comment >= 2x the previous median and reply-back share rises.	n/a

		followers in core topics; stop emoji/one-word replies. *Compare:* The account's previous 4 weeks of conversational comments Secondary: share of comments that receive a reply, new accounts engaging with the client's posts · Confounders: parent-post reach, parent author · Minimum period: 2-3 weeks			
3	**Lead-magnet delivery without mass identical replies** diagnostic: supported, low	Delivering lead magnets via DM/pinned reply instead of hundreds of identical public 'Sent' replies keeps conversions and removes a duplicate-reply risk. *Do:* For half of lead-magnet posts (alternating), deliver via one pinned reply + DM automation; vary wording. *Compare:* Alternating lead-magnet posts with the current 'Sent' replies Secondary: views vs baseline of the lead-magnet post, DMs delivered · Confounders: lead-magnet topic, CTA wording · Minimum period: 4 weeks (needs >= 4 lead-magnet posts per arm)	keyword comments per lead-magnet post	Switch if comments per post are within 15% of control.	n/a
4	**Evaluate, keep or drop**	n/a *Do:* Re-run the full audit pipeline; compare each experiment's primary metric against its decision rule. *Compare:* n/a Secondary: · Confounders: seasonality · Minimum period: end of week 4	decision per experiment	Retain experiments that met their rule; extend underpowered ones rather than calling them.	n/a

12. The fourteen questions

1. Why is the account receiving its current level of reach?

Median authored post: 1258 views (n=157). Ranked potential drivers: 1. Content format drives most of the reach difference (high); 2. Low conversation generation per impression (medium); 3. Off-core topics underperform core topics (medium); 4. Reach depends on rare breakouts (high); 5. Commenting is low-depth and concentrated on a small circle (medium) statistical association

2. Which content types perform best and worst?

Best: format=x_article 15.12x (n=30); narrative=tutorial 6.10x (n=46); topic=AI agents & GTM automation 4.18x (n=33); narrative=promotion 3.79x (n=10). Worst: narrative=other 0.21x (n=28); topic=personal & lifestyle 0.24x (n=9); topic=other 0.26x (n=15); topic=AI industry & tools news 0.27x (n=12). statistical association

3. Which posting behaviours are associated with better results?

length_bucket=short (15-49) 0.18x (n=24); length_bucket=very short (<15 words) 1.97x (n=56); weekday=Monday 2.43x (n=30); weekday=Saturday 0.27x (n=12); daypart_utc=06-11 0.20x (n=21); daypart_utc=12-17 1.91x (n=68). Same-day order (median vs baseline): 1st 1.10x (n=117), 2nd 1.38x (n=26), 3rd+ 0.23x (n=9) statistical association

4. Is the commenting strategy effective?

226 conversational comments in the window; 37% are substantive and 39% are generic, repetitive or humour. The top 5 accounts receive 35% of comments; median exposure is 47.0 views per comment. account observed

5. Which interactions are producing meaningful audience exposure?

Top comment targets by views: @vincent_alonzi (2 comments, median 309 views), @jacky_gng (2 comments, median 268 views), @512banque (4 comments, median 263 views), @antinertia (4 comments, median 223 views), @YannDecoopman (2 comments, median 195 views) account observed

6. Are there observable signs of potential visibility limitations?

0 of 157 mature authored posts have fewer than 100 views; reach is continuous with the account's normal range. No public signal indicates a visibility restriction; this audit cannot observe X's visibility-filtering decisions. insufficient

7. What does the published X algorithm actually tell us?

Ranking is a weighted sum of per-viewer predicted action probabilities (reply 5.0, like 0.66, share-via-copy-link 20.0, follow 4.0, profile click 0.0, report -234.0 in published defaults), with author-diversity decay within a feed, a 0.75 out-of-network factor, a 48-hour age cutoff and candidate retrieval driven by viewers' engagement history. See the findings registry. repo verified

8. Which common algorithm claims are unsupported?

TweepCred below 65 limits an account to 3 tweets being considered (not found in repo); Posts with external links get a 30-50% reach penalty (not found in repo); Using 3 or more hashtags gets a post penalised (not found in repo); Premium/verified accounts get a 4x in-network and 2x out-of-network boost (not found in repo); One report cancels out 468 likes (contradicted); Replies are worth 27x a like (contradicted); A reply the author engages with is worth 75x (or 150x) a like (not found in repo); Bookmarks are worth 20x a like (not found in repo); A good click / dwell (2+ minutes) is worth 10-11x a like (contradicted); Scores decay with a 6-hour half-life (contradicted) unsupported

9. How does the account compare with relevant peers?

authored per week: 7.58 vs peer median 13.86; views per 1k followers: 101.87 vs peer median 82.67; eng per 1k views: 9.25 vs peer median 22.73; replies per 1k: 0.93 vs peer median 1.83; bookmarks per 1k: 2.14 vs peer median 3.86 statistical association

10. What should the account stop doing?

Reduce: Length: short (15-49) (0.18x median, n=24); Reduce: Educational value (0-3): 0.0 (0.19x median, n=34); Reduce: Time of day (UTC): 06-11 (0.20x median, n=21); Reduce: Narrative structure: other (0.21x median, n=28); Reduce: Opening hook type: other (0.22x median, n=24); Stop publishing a 3rd authored post on the same day; Stop generic/emoji-only comments; they earn little exposure; Stop mass identical public 'Sent' replies (test DM/pinned delivery) statistical association

11. What should the account continue doing?

Keep: Content format: x_article (15.12x median, n=30); Keep: Opening hook type: how_to (8.87x median, n=33); Keep: Educational value (0-3): 2.0 (6.97x median, n=46); Keep: Opening hook type: story (6.24x median, n=8); Keep: Narrative structure: tutorial (6.10x median, n=46) statistical association

12. What should the account start doing?

Start measuring: daily metric snapshots + follower history (monitoring); Start the randomised hook and closing-prompt tests (see 30-day plan); Start substantive comments on larger in-niche accounts within their first hour; Start replying to early commenters within 60 minutes to build conversation depth hypothesis

13. Which changes should be tested first?

Week 1: Freeze the baseline and fix data gaps; Week 1: Space authored posts (max 2/day, >= 6h apart); Week 2: Article-anchored week vs mixed week; Week 2: Outcome-first hooks vs current openers; Week 2: Specific closing prompt + first-hour author replies hypothesis

14. How will we measure whether those changes work?

Each experiment has a primary metric and decision rule (see 30-day plan); judged on medians of views vs trailing baseline, replies/bookmarks per 1k views. Sample sizes assume a two-sided test at alpha 0.05 and 80% power on log views, using the residual spread after month/format/topic controls. Effects smaller than ~1.3x cannot be detected within 30 days at current posting rates. Needs ~393 posts per arm; at the current 6.3 authored posts/week that is ~126 weeks. hypothesis

Appendix — common algorithm claims checked against the repository

Claim	Verdict	Explanation
X uses a TweepCred-style reputation score for accounts	partially supported	There is no "TweepCred" in this repo, but user-cred-v2 computes a PageRank-based per-account score (0-100) from follow edges, with premium accounts and like/repost engagement in the teleport set. In this code the score is consumed by spam/abuse systems as an exemption threshold, not by the For You value model. user-cred-v2/UserCredV2.scala:9-20 botmaker-rules/scarecrow/derived-feature/IsHighPageRankUser.df:1-28 xai-value-model/inputs.rs:3-17
TweepCred below 65 limits an account to 3 tweets being considered	not found in repo	No TweepCred, no threshold of 65, and no per-author cap of 3 tweets were found in the published ranking code. The user-cred thresholds visible here are 54 (scarecrow high-PageRank), 50.0 (abuse-enforcement skip) and 66 (Grox reply-label exemption), and they gate spam labelling, not tweet counts. Absence here does not prove absence in production. abuse-enforcement-service/service-lib/rules/enforcement_user.yaml:32-38 grox/flows/reply_spam/constants.py:9-9
Posts with external links get a 30-50% reach penalty	not found in repo	The published weighted score has no link-penalty term; open_link is a positive head (weight 0.2). A has_url content feature exists only in an experimental debias payload that is off by default and is not read by the published value model. Link-specific effects could still arise from model predictions or URL-quality spam rules, which this does not rule out. home-mixer/params/param.rs:373-373 home-mixer/params/param.rs:502-507 home-mixer/models/content_features.rs:8-20
Using 3 or more hashtags gets a post penalised	not found in repo	No hashtag count or hashtag penalty appears in home-mixer, vm-ranker or xai-value-model; hashtags appear only in SimClusters entity extraction, muted keyword handling, ads tests and Agatha data loaders. This is absent from the public ranking code, not proven absent in production.
Premium/verified accounts get a 4x in-network and 2x out-of-network boost	not found in repo	No verification or premium multiplier appears in the For You scoring code; the only post-fusion multipliers are author diversity and the out-of-network factor. Premium status does appear in user-cred-v2 as the uniform teleport set, and gray/grey-badge status exempts accounts from some spam labels, which are indirect effects, not a ranking boost. xai-value-model/scoring.rs:12-26 user-cred-v2/UserCredV2App.scala:173-175
One report cancels out 468 likes	contradicted	The ratio comes from ReportWeight -234 vs a favorite weight of 0.5, but the weights multiply each viewer's predicted probabilities, not raw counts, and the code comment explicitly calls this reading incorrect. The current home-mixer favorite default is 0.66, so even the ratio differs by file. home-mixer/params/param.rs:320-330 home-mixer/params/param.rs:457-457
Replies are worth 27x a like	contradicted	Current defaults are ReplyWeight 5.0 and FavoriteWeight 0.66 in home-mixer (0.5 in vm-ranker), i.e. roughly 7.6x-10x, and they weight predicted probabilities rather than counts. The 27x figure may describe the 2023 heavy ranker, which is not in this repo. home-mixer/params/param.rs:345-346 vm-ranker/params.rs:3-4
A reply the author engages with is worth 75x (or 150x) a like	not found in repo	The current Phoenix heads have no "reply engaged by author" prediction; the weighted score sums the listed heads only. These numbers come from the 2023 heavy-ranker configuration, which is not part of this repository. xai-value-model/phoenix_scores.rs:1-31
Bookmarks are worth 20x a like	not found in repo	There is no bookmark prediction head or bookmark weight in the scoring code; bookmarks appear as viewer engagement-history signals and candidate counts. A weight of 20.0 exists for share_via_copy_link, which is a different action. xai-value-model/phoenix_scores.rs:1-31 home-mixer/params/param.rs:388-393 home-mixer/models/engagement_signals.rs:20-26
A good click / dwell (2+ minutes) is worth 10-11x a like	contradicted	Current defaults are ClickWeight 0.3, DwellWeight 0.05, ContDwellTimeWeight 0.006 and ContClickDwellTimeWeight 0.4 (per unit of predicted dwell time); there is no "good click" head with weight 10-11. The figure appears to come from the 2023 heavy ranker, not this code. home-mixer/params/param.rs:372-372 home-mixer/params/param.rs:432-437
Scores decay with a 6-hour half-life	contradicted	The published weighted score has no time-decay term; freshness is enforced by a hard 48-hour age cutoff. The only half-life found in home-mixer is 8 hours, used to normalise views in the popular-posts retrieval source, not to decay ranking scores. Recency could still be learned implicitly by the prediction model. xai-value-model/scoring.rs:69-140 home-mixer/params/config.rs:40-40 home-mixer/util/popular_posts.rs:36-46

The first 30 minutes decide a post's distribution	not found in repo	No 30-minute decision window exists in the ranking code. Related time windows are the cold-start lift for posts under 7,200 s old with fewer than 200 home impressions, a 0.5-hour age floor in popular-posts view normalisation, and the 48-hour age cutoff. Absence here is not proof of absence in production. home-mixer/params/param.rs:562-567 home-mixer/util/popular_posts.rs:36-46
An Author Diversity Scorer throttles accounts that post frequently	partially supported	An author-diversity multiplier exists, but it applies within a single viewer's candidate set: the k-th additional post from the same author is multiplied by 0.75 * 0.5^k + 0.25. It does not count how often an account posts overall. xai-value-model/scoring.rs:183-204 vm-ranker/params.rs:175-186
Grok reads every post and down-ranks negative sentiment	not found in repo	Grox runs LLM/VLM classifiers for safety, spam, reply spam and a "banger" screen, but no sentiment score or sentiment term appears in the ranking formula, and Grox flows are gated (e.g. multimodal embedding uses an adaptive minimum like count starting at 20,000). Prompts are withheld, so the exact criteria cannot be checked. grox/flows/ptos/prompts.py:11-11 grox/flows/mm_emb/adaptive_min_fav.py:6-19
Spelling errors cause a deboost	not found in repo	No spelling or typo check appears in the ranking or filtering code. Repo-wide regex matches are unrelated substrings (e.g. "tyPo" inside identifiers in a Grox file, a CLI-parsing test name in Phoenix). Not proof of absence in production; Grox prompts are withheld.
Mass unfollowing triggers a 3-month shadowban	not found in repo	No unfollow-based penalty or 3-month restriction appears in home-mixer, visibility-filtering or the published enforcement rules. bdsm's FollowBot head targets automated follow behaviour, but its thresholds are redacted and no duration like 3 months is published. bdsm/runtime/sink_policy.yaml:15-20
SimClusters has about 145k communities	partially supported	The SimClusters code refers to a model version named Model20M145K2020, which is consistent with ~145K clusters, but no explicit cluster count constant was found. SimClusters is used in home-mixer as one retrieval source. simclusters/simclusters_v2/summingbird/common/SimClustersProfile.scala:145-145
Video posts get 10x reach	not found in repo	No video multiplier exists in the scoring code. Video heads carry small weights by default (video_open 0.02 in home-mixer, video quality view 0.0). Video can still affect predicted engagement and has dedicated carousel and immersive surfaces, so effects on reach are not ruled out. home-mixer/params/param.rs:366-371 home-mixer/params/param.rs:380-380
Replies between mutual followers get boosted	partially supported	A bidirectional-follow boost exists (default +15 on the reply weight), but it applies to original posts from authors the viewer mutually follows, raising the weight on the predicted probability that the viewer replies. Replies and reposts are explicitly excluded from eligibility. xai-value-model/inputs.rs:19-23 home-mixer/params/param.rs:347-352 docs/BIDIRECTIONAL_BOOST_CHANGE.md:5-5
Profile clicks are a major ranking signal	contradicted	ProfileClickWeight defaults to 0.0 and profile-visit seconds are weighted 0.0 (home-mixer hard-codes it to 0.0), so predicted profile clicks do not add to the published weighted score at defaults. Experiments could set non-zero values. home-mixer/params/param.rs:374-379 home-mixer/scorers/value_model.rs:38-40
Out-of-network posts are down-weighted	supported	Out-of-network candidates (and, by default, in-network replies and reposts) are multiplied by effective_oon_weight, which defaults to OonWeightFactor 0.75 (0.5 for topic requests). Out-of-network posts also face extra visibility drop rules. xai-value-model/scoring.rs:206-215 vm-ranker/params.rs:211-216
Posting too often cannibalises your own reach	partially supported	Within one viewer's feed, repeated posts by the same author are discounted (author diversity) and only one post per conversation is kept, so several simultaneous posts compete with each other for that viewer. Nothing in the code penalises posting frequency globally or across viewers. xai-value-model/scoring.rs:183-185 home-mixer/filters/dedup_conversation_filter.rs:19-36
Mass blocking/reporting by bad actors suppresses an account's reach	unverifiable	For ranking, the code comment states that weights apply to each viewer's predicted probabilities and that only actions on posts served in Home Timeline count, so coordinated reports mainly affect similar users. However, Agatha computes account-level blocks-per-fav and reports-per-fav rates that feed account models, and the thresholds turning those into labels are not published. home-mixer/params/param.rs:332-344 agatha/scalding/labels/rate_based_labels/RateBasedLabels.scala:80-90